

Dynamic Causal Effects in a Nonlinear World: the Good, the Bad, and the Ugly*

Michal Kolesár
Princeton University

Mikkel Plagborg-Møller
Princeton University

July 15, 2025

Abstract: Applied macroeconomists frequently use impulse response estimators motivated by linear models. We study whether the estimands of such procedures have a causal interpretation when the data generating process is in fact nonlinear. We show that vector autoregressions and linear local projections onto observed shocks or proxies identify weighted averages of causal effects regardless of the extent of nonlinearities. By contrast, identification approaches that exploit heteroskedasticity or non-Gaussianity of latent shocks are highly sensitive to departures from linearity. Our analysis is based on new results on the identification of marginal treatment effects through weighted regressions, which may also be of interest to researchers outside macroeconomics.

Keywords: dynamic treatment effect, impulse response, local projection, semiparametric identification, structural vector autoregression.

1 Introduction

Impulse response functions are key objects in macroeconomic analysis. Since they measure dynamic causal effects of surprise changes in policy or fundamentals on subsequent macroe-

*Email: mcolesar@princeton.edu, mikkelpm@princeton.edu. We are grateful for comments from Joshua Angrist, Paul Bousquet, Alessandro Casini, Runtong Ding, Florian Gunsilius, Peter Hull, Arthur Lewbel, Daniel Lewis, Dexin Li, Adam McCloskey, Geert Mesters, Aureo de Paula, Julius Schäper, Christian Wolf, participants at the JBES session at the 2025 ASSA meetings and the Federal Reserve Bank of New York seminar, two anonymous referees, the editor Atsushi Inoue, and particularly from our discussants: Sílvia Gonçalves, Ed Herbst, Ana María Herrera, Ben Johannsen, Òscar Jordà, Toru Kitagawa, Elena Pesavento, Weining Wang, and Mengshan Xu. Eric Qian and Nelson Matthew Tan provided excellent research assistance. Kolesár acknowledges support by the National Science Foundation under Grant SES-2049356. Plagborg-Møller acknowledges support from the National Science Foundation under Grant SES-2238049 and from the Alfred P. Sloan Foundation.

conomic outcomes, they provide calibration targets for structural modeling and help validate model predictions. They also inform optimal economic policy questions, both directly and indirectly (Christiano, Eichenbaum, and Evans, 1999; McKay and Wolf, 2023).

Applied researchers typically report impulse response estimators motivated by linear time series models, such as vector autoregression (VARs) or local projections. Although there exists a wealth of nonlinear alternatives (Fan and Yao, 2003; Herbst and Schorfheide, 2016; Kilian and Lütkepohl, 2017, Chapter 18), linear methods are attractive due to their simplicity and the difficulty of clearly detecting nonlinear relationships in typical macroeconomic data. At the same time, both macroeconomic theorists and policymakers think nonlinearities are important: structural models with essential nonlinearities have become dominant in recent decades, and many economic policy debates concern state-dependence and asymmetries. How can we justify using linear methods if we think the world is a nonlinear place?

This paper studies the causal interpretation of impulse response estimators based on linear models when the data is generated by an essentially unrestricted nonparametric structural model. We first deliver *good* news for linear local projection or VAR estimators that project directly on an observed shock or proxy: their estimand (i.e., probability limit) equals a weighted average of the true nonlinear causal effects, regardless of the extent of nonlinearities in the data generating process (DGP). By contrast, the news is *bad* or even *ugly* for estimators that identify latent shocks via heteroskedasticity or non-Gaussianity: they generally do not estimate a meaningful causal summary under departures from linearity. Thus, the hard work needed to directly measure shocks (or proxies) using historical or institutional data buys insurance against nonlinearities that other identification approaches lack.

Our good news are based on an extension of the results in Yitzhaki (1996) and Rambachan and Shephard (2021): impulse response estimands from linear local projections and VARs that project on observed shocks or proxies correspond to positively-weighted averages of marginal effects—causal effects of changing the shock variable from x to $x + \delta$ for infinitesimally small δ , averaged out over all (past, present, and future) shocks other than the contemporaneous shock of interest, and weighted over different baseline values x for the shock variable. Thus, these estimands provide a scalar causal summary of the full richness of the nonlinear causal effects; the positive weights ensure that the researcher gets the sign right if the true marginal effects are uniformly positive or negative. Our assumptions drop restrictions imposed in the existing literature that ruled out models with kinks or discontinuous regime-switches or shocks with unbounded support.

In a nonlinear DGP, both the sign and the magnitude of the causal effects can depend

on the baseline shock value (e.g., whether it is positive or negative), so how these values are weighted can matter a lot. Fortunately, as we illustrate using several empirical examples, the weight function used by local projections and VARs is straightforward to estimate and report. In many applications, the researcher does not directly observe the shock of interest but only a proxy, also known as an external instrument (Stock and Watson, 2018). In this case, we show that an easily-interpretable monotonicity condition is required to guarantee a positive weight function. However, we also show that when control variables are needed to isolate a true shock (i.e., recursive or Cholesky identification), then positive weights can only be guaranteed if the linearly residualized shock is *nonlinearly* unpredictable by the controls, which may be a strong assumption in the absence of detailed institutional knowledge and high-quality data.

One implication of these results is that linearity-based estimators are useful even when economic theory predicts a nonlinear relationship between the shock and the outcome of interest. For example, if the outcome variable has limited support, such as when it is binary or censored (say, due to a zero lower bound), nonlinearities are inherently present. If one is interested in *characterizing* the nonlinearities, then it makes sense to model them, and it is of course always a good idea to plot the raw data regardless. However, if one is interested in an overall summary of *marginal effects*, then linear local projections and VARs are theoretically coherent estimators, as discussed earlier. In fact, we show that directly modeling nonlinearities can be counterproductive unless the researcher is confident in their modeling: under functional form misspecification, local projections with higher-order terms still estimate a weighted average of marginal effects, but some of the weights may be negative, which risks getting the sign of the causal effects wrong. This echoes the message from an earlier JBES lecture by Angrist (2001) that in a cross-section context with limited dependent variables, linear methods provide more robust estimates of treatment effects than nonlinear ones.

When there is a dearth of direct shock measures or proxies, applied researchers frequently resort to identification via heteroskedasticity (Sentana and Fiorentini, 2001; Rigobon, 2003; Lewbel, 2012). Unfortunately, we show that these estimation approaches are sensitive to the assumption that the structural model is linear: the estimand can easily be nonzero when there is no causal effect, or negative when the true shock has a uniformly positive effect on the outcome of interest. Fixing these issues while still delivering informative inference appears difficult, since a natural nonparametric generalization of the identification strategy yields very wide identified sets. The intuition for these negative results is that the identification

exploits a source of exogenous variation that shifts the *scale* of the latent shock of interest but not its mean. Without strong functional form assumptions, this type of exogenous variation is uninformative about the effect of a *location* shift in the shock on the conditional mean of the outcome, i.e., the impulse response. However, a silver lining is that the linear model delivers testable restrictions.

The sensitivity to nonlinearity is even greater for identification via non-Gaussianity (Comon, 1994; Gouriéroux, Monfort, and Renne, 2017; Lanne, Meitz, and Saikkonen, 2017). Also known as independent components analysis (ICA), this identification approach has recently increased significantly in popularity in the VAR literature. We show that the non-parametric analogue of the identification assumptions yields an identified set so large that effectively any function of the data can be construed as a “shock”. Intuitively, the mere assumptions that the latent shocks are independent and non-Gaussian are vacuous in a non-parametric context: any collection of random variables can always be represented as some nonlinear function of independent uniformly distributed random variables. Moreover, we give examples of simple DGPs featuring slight nonlinearity for which *any* linearity-based ICA procedure is highly biased asymptotically, yet in these DGPs one cannot reject the validity of the linear model.

The building block underlying most of the above findings is a set of results on the identification of weighted averages of marginal treatment effects using weighted regressions, which connects our analysis to a large literature in microeconometrics (e.g., Yitzhaki, 1996; Newey and Stoker, 1993; Angrist and Krueger, 1999; Goldsmith-Pinkham, Hull, and Kolesár, 2024). We extend existing results in this literature by unifying the treatment of continuous, discrete, and mixed regressors, and by substantially weakening the regularity conditions: we allow for regressors with unbounded support, impose minimal regularity on the regression function, and our moment conditions essentially only require the existence of the probability limit of the regression estimator.

An important limitation of our results is that they only concern identification. While we are motivated by the observation that full-fledged nonparametric estimation is challenging in realistic macroeconomic data sets, we do not explicitly analyze the precision or small-sample bias of the estimators we study. We refer to Herbst and Johannsen (2024) for a discussion of finite-sample biases of local projections and VARs in linear models. Another limitation is that we do not consider identification via non-recursive short-run restrictions, long-run restrictions, or sign restrictions.

LITERATURE. Pioneering work on semiparametric causal time series analysis includes Gallant, Rossi, and Tauchen (1993), Potter (2000), White (2006), White and Kennedy (2009), Angrist and Kuersteiner (2011), and Angrist, Jordà, and Kuersteiner (2018), see also Gonçalves, Herrera, Kilian, and Pesavento (2021, 2024a,b), Gouriéroux and Lee (2023), and Kitagawa, Wang, and Xu (2023) for recent contributions. Our result on the causal interpretation of local projections with observed shocks is very closely related to Rambachan and Shephard (2021) and subsequent work by Caravello and Martínez Bruera (2024) and Casini and McCloskey (2024).

As for identification via heteroskedasticity or non-Gaussianity, we are not aware of other work in a nonparametric vein. Montiel Olea, Plagborg-Møller, and Qian (2022) criticize linearity-based versions of these identification strategies for being seemingly sensitive to functional form assumptions, and potentially being subject to weak identification. The present analysis quantifies this sensitivity more precisely by deriving both the identified sets for the nonparametric analogues of these identification assumptions, and the estimands of linearity-based procedures.

OUTLINE. Section 2 defines a nonparametric framework for identification of dynamic causal effects. Section 3 argues that local projection and VAR estimands based on observed shocks or proxies have a robust causal interpretation regardless of the extent of nonlinearities. Sections 4 and 5 show, on the other hand, that estimands based on identification through heteroskedasticity or non-Gaussianity are sensitive to the assumption that the structural function is linear. Section 6 provides the theoretical basis for the results in the earlier parts of the paper by extending results from the microeconomic literature on the interpretation of regression estimators as weighted marginal treatment effects; this section may be of independent interest for readers outside macroeconomics. Section 7 concludes. Technical details and proofs are relegated to the online supplement.

2 Nonparametric framework for dynamic causality

In this section we set up a nonparametric framework for dynamic causal identification.

2.1 Model

We are interested in the dynamic response of a scalar outcome variable Y_t to an impulse in the scalar shock variable X_t . As a leading example, one may think of X_t as a variable

controlled by a policy-maker, such as a surprise change in the policy interest rate set by the central bank. For ease of exposition, we restrict attention to continuously distributed shocks X_t for now, but [Section 6](#) shows that our results generalize to handle continuous, discrete, or mixed distributions in a unified manner.

The outcome variable is determined by an underlying dynamic structural model. Our causal framework doesn't restrict this model; we only assume that when evaluated h periods after the realization of the shock X_t , the outcome admits the nonparametric structural representation

$$Y_{t+h} = \psi_h(X_t, \mathbf{U}_{h,t+h}) \quad \text{for all } t, h \geq 0. \quad (1)$$

For each horizon h , $\psi_h(\cdot, \cdot)$ is an unknown measurable function that we call the *structural function*, while $\mathbf{U}_{h,t+h}$ is a vector of all variables (dated before, on, and after time t) that causally affect Y_{t+h} , other than X_t . Without restrictions on $\mathbf{U}_{h,t+h}$ or the structural function, the representation (1) is without loss of generality. In typical recursive time-series models, however, $\mathbf{U}_{h,t+h}$ will contain the vector \mathbf{Y}_{t-1} of observed data at time $t-1$ as well as shocks dated $t, t+1, \dots, t+h$, but exclude X_t and shocks dated after $t+h$ ([White, 2006](#); [White and Kennedy, 2009](#); [Caravello and Martínez Bruera, 2024](#); [Gonçalves, Herrera, Kilian, and Pesavento, 2024b](#)). A leading special case is the linear structural VAR model, which additionally implies that ψ_h is linear in both X_t and $\mathbf{U}_{h,t+h}$ (e.g., [Kilian and Lütkepohl, 2017](#), Chapter 4.1).

We assume throughout that it is meaningful to think of varying X_t while keeping $\mathbf{U}_{h,t+h}$ constant, so that the random function $Y_{t+h}(x) \equiv \psi_h(x, \mathbf{U}_{h,t+h})$ defines a *potential outcome function* at horizon h . [Angrist and Kuersteiner \(2011\)](#), [Angrist, Jordà, and Kuersteiner \(2018\)](#), and [Rambachan and Shephard \(2021\)](#) work with this potential outcome notation $Y_{t+h}(x)$, and keep all other past, present, and future shocks (captured by $\mathbf{U}_{h,t+h}$ in our model) implicit. Our structural function framework (1) is mathematically equivalent, but facilitates comparisons with the linear structural VAR literature.

For now, we focus on the case where X_t is a “shock” (such as a surprise change in a policy instrument), and assume that it is independent of $\mathbf{U}_{h,t+h}$:

$$X_t \perp\!\!\!\perp \mathbf{U}_{h,t+h}. \quad (2)$$

This assumption is common in the literature, and essentially just normalizes the structural function ψ_h , so that its first argument captures the total causal effect of the shock X_t on Y_{t+h} , including its direct effect and any indirect effects, both contemporaneous and dynamic.

This is illustrated in the following simple example.

Example 1. Consider a univariate AR(1) model with endogenous regime switching:

$$Y_t = \rho_t Y_{t-1} + \tau \varepsilon_t + \nu_t,$$

with regime-dependent parameter $\rho_t = \rho_1 S_t + \rho_0(1 - S_t)$ and binary regime $S_t = \mathbb{1}\{\varepsilon_{t-1} + \xi_{t-1} \leq 0\}$, and where ρ_0, ρ_1, τ are constants. Assume that ε_t, ν_t , and ξ_t are i.i.d. and mutually independent, and that we observe the shock $X_t = \varepsilon_t$. We can cast this model into the form required by [equations \(1\) and \(2\)](#) as follows. Define $\mathbf{U}_{h,t+h} \equiv (Y_{t-1}, S_t, \nu_t, \dots, \nu_{t+h}, \xi_t, \dots, \xi_{t+h-1}, \varepsilon_{t+1}, \dots, \varepsilon_{t+h})'$ and $\rho(\vartheta) \equiv \rho_1 \mathbb{1}\{\vartheta \leq 0\} + \rho_0 \mathbb{1}\{\vartheta > 0\}$ for $\vartheta \in \mathbb{R}$. Then, for all $h \geq 1$,

$$\begin{aligned} \psi_h(x, \mathbf{u}) = & \left\{ y_{-1}(\rho_1 s + \rho_0(1 - s)) + (\tau x + \nu) \right\} \rho(x + \xi) \prod_{\ell=1}^{h-1} \rho(\varepsilon_{+\ell} + \xi_{+\ell}) \\ & + \sum_{\ell=1}^h (\tau \varepsilon_{+\ell} + \nu_{+\ell}) \prod_{b=\ell}^{h-1} \rho(\varepsilon_{+b} + \xi_{+b}), \end{aligned}$$

where we have partitioned $\mathbf{u} = (y_{-1}, s, \nu, \nu_{+1}, \dots, \nu_{+h}, \xi, \xi_{+1}, \dots, \xi_{+(h-1)}, \varepsilon_{+1}, \dots, \varepsilon_{+h})'$. Notice that the function ψ_h captures the full dynamic effect of the shock variable $X_t = \varepsilon_t$: both the direct impact effect of ε_t on Y_t (which feeds forward to future periods), and the indirect, nonlinear effect of ε_t on the next-period regime S_{t+1} (which also feeds forward). ■

When X_t is not a shock but corresponds to a policy instrument (such as the policy interest rate) or any other kind of serially correlated variable, it will typically be correlated with $\mathbf{U}_{h,t+h}$, both because of the serial correlation and because it may be correlated with other determinants of Y_{t+h} —we discuss this case in [Section 3.3](#).

2.2 Causal effects

A familiar issue in nonlinear models is that there are multiple possible definitions of an impulse response, i.e., a dynamic causal effect. In a linear model, the effect of exogenously changing X_t from x_0 to x_1 is a linear function of the difference $x_1 - x_0$: it equals $\psi_h(x_1, \mathbf{U}_{h,t+h}) - \psi_h(x_0, \mathbf{U}_{h,t+h}) = \theta_h(x_1 - x_0)$ for some constant scalar θ_h . By contrast, in a nonlinear model, the effect is in general a nonlinear function of both x_1 and x_0 (not just their difference), and it also depends on the past history and the current and future nuisance shocks via $\mathbf{U}_{h,t+h}$.

In theoretical macroeconomic modeling, researchers often report the impulse responses with respect to a so-called “MIT shock”, which starts the economy at steady state, then hits the economy with a one-off impulse to X_t , and subsequently sets all other current and future shocks to zero: $\psi_h(X_t, \mathbf{0}) - \psi_h(0, \mathbf{0})$, where we normalize the steady-state values of X_t and $\mathbf{U}_{h,t+h}$ to zero. While computationally convenient, this impulse response concept has no empirical counterpart; moreover, for the measure to be policy relevant, it is necessary that the model satisfy certainty equivalence.

We shall instead focus on impulse responses (causal effects) defined as expected counterfactual changes in the outcome of interest, averaging out over all other shocks. Specifically, define the *average structural function*

$$\Psi_h(x) \equiv E[\psi_h(x, \mathbf{U}_{h,t+h})], \quad x \in \mathbb{R},$$

which corresponds to the expected potential outcome function. Here the expectation is taken over the marginal distribution of $\mathbf{U}_{h,t+h}$. The expectation is implicitly assumed to exist for all x . The average structural function measures the counterfactual average value of the future outcome Y_{t+h} that we would observe if the policy-maker engineered a particular fixed value x for the policy variable at time t , averaging out over the randomness caused by all other factors that influence the outcome independently of the policy decision at time t . Even though in some nonlinear models the structural function ψ_h is discontinuous in the policy variable x , the average structural function Ψ_h will typically be a smoother function of x , as it averages out over the realizations of other shocks. For example, this is the case in the regime-switching model in [Example 1](#) if ξ_t is continuously distributed. A certain amount of smoothness in Ψ_h will be important for the identification of causal effects, as we discuss below.

Typical data samples in macroeconomics are too small to permit accurate nonparametric estimation of the entire average structural function $x \mapsto \Psi_h(x)$. A pragmatic alternative is to target weighted averages of the structural function—average causal effects—or its derivatives—average marginal effects ([Rambachan and Shephard, 2021](#); [Gonçalves, Herrera, Kilian, and Pesavento, 2021, 2024b](#)). This paper focuses on estimation of *average marginal effects*

$$\theta_h(\omega) \equiv \int \omega(x) \Psi'_h(x) dx, \tag{3}$$

where $\omega(\cdot)$ is a weight function averaging across the baseline values of the shock variable X_t . We reserve the term *average marginal effect* to weight functions that are convex, i.e.,

$\omega(x)$ is nonnegative for all x and integrates to one, $\int \omega(x) dx = 1$. This ensures that $\theta_h(\omega)$ is a meaningful causal summary of the average structural function $\Psi_h(x)$ in that it prevents what [Small, Tan, Ramsahai, Lorch, and Brookhart \(2017\)](#) call a sign-reversal: if $\Psi'_h(x)$ has the same sign for all x (+, 0 or -), then $\theta_h(x)$ will also have this sign.¹ This property is particularly useful when qualitatively validating predictions of structural macroeconomic models.

Depending on the form of the weight function, $\theta_h(\omega)$ has two interpretations in terms of an *average causal effect* of a shock with magnitude $\delta > 0$,

$$\theta_h(\delta, \omega_0) \equiv \frac{1}{\delta} \int \omega_0(x) \{\Psi_h(x + \delta) - \Psi_h(x)\} dx. \quad (4)$$

First, the average marginal effect corresponds to the average causal effect for infinitesimally small shocks: $\theta_h(\omega) = \lim_{\delta \rightarrow 0} \theta_h(\delta, \omega)$, provided we can pass the limit as $\delta \rightarrow 0$ under the integral sign in (4). Second, if the weighting in (3) admits the integral representation $\omega(x) = \frac{1}{\delta} \int_{x-\delta}^x \omega_0(x) dx$, substituting $\Psi_h(x + \delta) - \Psi_h(x) = \int_x^{x+\delta} \Psi'(\chi) d\chi$ into (4) and changing the order of integration yields $\theta_h(\omega) = \theta_h(\delta, \omega_0)$. For this reason, focusing on average marginal effects is without loss of generality.

In a linear model, the weighting does not matter, since $\Psi'_h(x)$ does not depend on x . But in nonlinear models, it could matter greatly whether we attach most weight to positive or negative shocks, or to shocks with small or large magnitude. Therefore, accounting for the form of the weighting ω is important when using estimates of $\theta_h(\omega)$ to calibrate or validate structural macroeconomic models. In the next section, we discuss identification approaches that deliver weighted averages of marginal effects under a particular weighting scheme that depends on the shock distribution. In [Section 6](#), we discuss estimation approaches that target any pre-specified weighting scheme.

3 The good: observed shocks and proxies

If the researcher directly observes the shock of interest, or at least a valid proxy for it, then there is *good* news: conventional local projections or structural VAR impulse responses estimate average marginal effects with an interpretable weighting scheme, regardless of how

¹[Blandhol, Bonney, Mogstad, and Torgovitsky \(2022\)](#) call estimands with convex ω “weakly causal”. Convex weighting schemes also satisfy what [Robins, Sued, Lei-Gomez, and Rotnitzky \(2007\)](#) call “boundedness”: $\theta_h(\omega)$ lies in the support of $\Psi'_h(x)$.

nonlinear the underlying DGP is. Moreover, the weights can be estimated from the data, and we give several empirical examples of how to interpret them. In contrast to linear estimators, we demonstrate using a simple example that nonlinear extensions of local projections or VARs do not generally provide meaningful causal summaries under misspecification. Finally, we extend the analysis to shocks that are recursively identified, i.e., by controlling for covariates.

3.1 Identification with observed shocks

We start off by assuming that the researcher directly observes (or consistently estimates) the shock X_t of interest. This would be the case, for example, if the shock is identified through a “narrative approach”. See [Ramey \(2016\)](#) for several empirical examples.

Under the nonlinear structural model (1) and the shock independence assumption (2), the conditional expectation of the outcome given the shock,

$$g_h(x) \equiv E[Y_{t+h} \mid X_t = x], \quad (5)$$

nonparametrically identifies the average structural function:

$$\Psi_h(x) = E[\psi_h(x, \mathbf{U}_{h,t+h})] = E[\psi_h(x, \mathbf{U}_{h,t+h}) \mid X_t = x] = g_h(x). \quad (6)$$

Hence, in principle, we could estimate any weighted causal effect of interest by running a nonparametric regression of Y_{t+h} on X_t to obtain $g_h(\cdot)$ in the first step, and then averaging this function according to the desired weighting scheme in the second step, as suggested by [Gonçalves, Herrera, Kilian, and Pesavento \(2024a,b\)](#). In [Section 6](#), we discuss a complementary strategy that identifies the same estimand via weighted averages of the observed outcomes, and how both strategies can be combined. However, as discussed in more detail in [Section 6](#), these strategies may yield noisy and sensitive estimates in the relatively small samples available in macroeconomics.

INTERPRETATION OF LINEAR PROJECTION ESTIMATES. We take a cue from [Ram-bachan and Shephard \(2021\)](#) and instead aim for a less ambitious goal. Rather than targeting a pre-specified weighted average of causal or marginal effects, we focus on simple local projection and VAR estimators, which are relatively precise even with small sample sizes. We demonstrate that these simple estimators have an attractive robustness property: even though they are motivated by a linear model, when the DGP is nonlinear, their estimand can still be interpreted as an average marginal effect with a particular weight function.

The local projection estimator of [Jordà \(2005\)](#) estimates the impulse response of Y_t with respect to X_t at horizon h as the coefficient $\hat{\beta}_h$ in the ordinary least squares (OLS) regression

$$Y_{t+h} = \hat{\beta}_h X_t + \hat{\gamma}_h' \mathbf{W}_t + \text{residual}_{h,t+h}, \quad (7)$$

where \mathbf{W}_t is a vector of control variables (typically including a constant and lagged outcomes and shocks). For now, we will assume that the shock X_t is in fact a “shock”, so that it is linearly unpredictable using the controls: $\text{Cov}(X_t, \mathbf{W}_t) = 0$. Then the set of controls \mathbf{W}_t affects only the precision of $\hat{\beta}_h$, but not its probability limit. In particular, under standard stationarity and ergodicity assumptions, the local projection estimator $\hat{\beta}_h$ will converge in probability to the population projection coefficient

$$\beta_h \equiv \frac{\text{Cov}(g_h(X_t), X_t)}{\text{Var}(X_t)}. \quad (8)$$

[Plagborg-Møller and Wolf \(2021\)](#), Propositions 1 and 2) show that a VAR which includes X_t ordered first has the exact same population estimand (8), provided that the number of lags in the VAR is sufficiently large.² It is a textbook result that the linear function $\beta_h x$ provides the best *linear* approximation to the potentially nonlinear average structural function $g_h(x) = \Psi_h(x)$ (e.g., [Angrist and Pischke, 2009](#), Theorem 3.1.6), so that it approximates the average structural function in a *prediction* sense. However, this result is not directly informative about whether β_h has a *causal* interpretation—whether it can be interpreted as an average marginal effect if $\Psi_h(x)$ is nonlinear.

The following proposition shows that the local projection and VAR estimand (8) achieves our goal: it has a causal interpretation as an average marginal effect (3). The result is not new—it appeared previously in [Yitzhaki \(1996\)](#) and [Rambachan and Shephard \(2021\)](#); as we discuss below, the novelty lies in substantively weakening the regularity conditions.

Proposition 1. *Assume that X_t is continuously distributed on an interval $I \subseteq \mathbb{R}$ (the interval may be unbounded, and could equal \mathbb{R}), with positive and finite variance. Assume that the conditional mean g_h defined in (5) is locally absolutely continuous on I .³ Suppose finally that $E[|g_h(X_t)|(1 + |X_t|)] < \infty$ and $\int_I \omega_X(x) |g_h'(x)| dx < \infty$, where*

$$\omega_X(x) \equiv \frac{\text{Cov}(\mathbb{1}\{X_t \geq x\}, X_t)}{\text{Var}(X_t)}. \quad (9)$$

²If X_t is linearly unpredictable from lagged data, it is sufficient that the lag length weakly exceed h .

³That is, absolutely continuous on any compact interval contained in I .

Then the estimand (8) satisfies

$$\beta_h = \int_I \omega_X(x) g'_h(x) dx,$$

and the weight function ω_X has the following properties:

- (i) It is convex: $\omega_X(x)$ is non-negative for all x , and integrates to one, $\int_I \omega_X(x) dx = 1$.
- (ii) It is hump-shaped: monotonically increasing from 0 to its maximum for $x \leq E[X_t]$, and then monotonically decreasing back to 0 for $x \geq E[X_t]$.
- (iii) It depends only on the marginal distribution of X_t , and not on the conditional distribution of Y_{t+h} given X_t .

Combined with the identification result (6) for the average marginal effect, **Proposition 1** shows that linear local projections and VARs remain useful in a nonlinear world: they estimate an average causal effect $\theta_h(\omega_X) = \int \omega_X(x) \Psi'_h(x) dx$ for infinitesimal shocks, with a convex weighting scheme ω_X . Furthermore, the scheme gives most weight to shocks close to the mean $E[X_t]$, with little weight given to extreme values. In the special case where X_t is normally distributed, **Proposition 1** reduces to Stein's lemma (Lemma 1 in **Stein, 1981**): the weight function ω_X reduces to the normal density function, so that β_h equals the expected marginal effect, $E[\Psi'_h(X_t)]$, as noted by **Yitzhaki (1996)**. The fact that the weighting scheme depends only on the marginal distribution of X_t and not the particular outcome variable Y_{t+h} or horizon h allows for comparisons of average marginal effects for different outcomes or across different horizons h .⁴ If the true DGP is in fact linear, then the weighting of course does not matter, and we recover the conventional linear impulse response.

While we focus here on interpreting the proposition in the context of the causal model in **Section 2**, the result does not require the structural assumptions (1)–(2). This is relevant in settings in which the conditional mean $g_h(x) = E[Y_{t+h} | X_t = x]$ is a useful descriptive object even if it does not have a direct causal interpretation.

The assumption that X_t is continuously distributed can be dropped without changing the result, as we show in **Section 6**. In cases where there are gaps in the support of X_t , such as when the shock is discrete or mixed, one just needs to extend the definition of the conditional mean function g_h to the whole interval I by linear interpolation. To our

⁴Since the weighting does not depend on the outcome variable or horizon, multipliers (i.e., ratios of cumulative impulse response estimands, see **Jordà, 2023**, Section 3) can be expressed as a ratio of two average marginal effects with the same weight function, where we now additionally average across horizons.

knowledge, this unification of the treatment of continuous, discrete, and mixed distributions is novel.

Even in the case of a continuously distributed shock, the assumptions in [Proposition 1](#) are substantively weaker than those in the literature, and accommodate all textbook linear models as well as a wide range of nonlinear models. The assumption that g_h is locally absolutely continuous is necessary to ensure that weighted marginal effects are well-defined. As discussed earlier, this assumption will typically hold even in models with discrete regimes or kinks, since the expectation (6) averages over the distribution of the nuisance shocks. The moment conditions and integrability condition $\int_I \omega_X(x) |g'_h(x)| dx < \infty$ just ensure that the estimand β_h and the weighted marginal effect exist.⁵ In contrast, the original work by [Yitzhaki \(1996\)](#) does not provide a formal proof or regularity conditions on g_h (neither does the discussion by [Angrist and Pischke, 2009](#), pp. 78 and 110). Analogous results in [Ram-bachan and Shephard \(2021\)](#), [Graham and de Xavier Pinto \(2022\)](#), [Caravello and Martínez Bruera \(2024\)](#), and [Casini and McCloskey \(2024\)](#) require the potential outcome function (not its expectation) to be smooth, which rules out models with kinks or discrete regimes, and require the interval I to be bounded, which rules out the textbook case of normally distributed shocks.⁶ The restrictiveness of these conditions led [Gonçalves, Herrera, Kilian, and Pesavento \(2024b, Appendix C\)](#) to question the applied relevance of the causal interpretation of the estimand (8), but our weaker conditions demonstrate that this concern is unfounded when interest centers on average marginal effects.

ESTIMATING THE WEIGHT FUNCTION. As argued by [Angrist and Krueger \(1999\)](#) for the case of discrete X_t , the weight function ω_X defined in (9) can be estimated in the data. This allows the researcher to gauge which weighted causal effect is being estimated: does it attach most weight to negative or positive shocks, small or large shocks? Since the weight function depends only on the shock variable itself and not the outcome variable or the impulse response horizon, it is only necessary to estimate a single function. We therefore recommend that researchers always estimate and plot this function.

Estimation is simple: $\omega_X(x)$ equals the slope coefficient in a (population) regression of the indicator $\mathbb{1}\{X_t \geq x\}$ on X_t . This regression can be implemented in the data via OLS,

⁵Lemma B.1 in Supplemental Appendix B shows that for the integrability condition to hold, it is sufficient to assume the tails of $g_h(x)$ are monotone.

⁶Our result is not strictly more general than that of [Casini and McCloskey \(2024\)](#), since they allow for non-stationary potential outcomes.

separately for each value x of X_t observed in the data.⁷ In applications, it may also be of interest to report an integral $\int_{\underline{x}}^{\bar{x}} \omega_X(x) dx$ of the weight function over an interval $x \in [\underline{x}, \bar{x}]$. Supplemental Appendix A.1 shows that we can estimate this integral by the slope coefficient in an OLS regression of $M_t \equiv \max\{\min\{X_t, \bar{x}\}, \underline{x}\}$ on X_t . In particular, to estimate the total weight $\int_0^\infty \omega_X(x) dx$ given to positive shocks, we simply regress $M_t \equiv \max\{X_t, 0\}$ on X_t .

To illustrate, we now empirically estimate the weight function ω_X for various macroeconomic shocks considered in the handbook chapter by [Ramey \(2016\)](#). We use [Ramey's](#) replication code and data off the shelf. In particular, prior to computing weights, all shocks are residualized on the same control variables that she uses in her VARs and local projections. The estimates of the weight functions are obtained from OLS regression output, as described above. To demonstrate the ease of implementation, all steps of the computations are carried out in Stata, like [Ramey's](#) replication code.⁸

[Figure 1](#) shows the estimated weight functions for four identified government spending shocks from the applied literature. Note that the shocks are not entirely comparable due to differences in their precise definitions and sample periods. The [Blanchard and Perotti \(2002\)](#) and [Fisher and Peters \(2010\)](#) shocks, which are intended to capture general government spending shocks, yield approximately symmetric weight functions. By contrast, the [Ben Zeev and Pappa \(2017\)](#) and [Ramey \(2011\)](#) shocks, which capture news about future defense spending, generate weight functions that are skewed towards positive shocks. In fact, both these shocks exhibit a large positive outlier in the 3rd quarter of 1950 (the onset of the Korean War), reflected in the fat right tail of the weight functions. In other words, impulse responses from local projections or VARs estimated off the latter two shocks will largely reflect the causal effects of sharp military *buildups*, rather than retrenchments. This is important to remember when using empirical impulse responses to discipline structural models that feature asymmetries (such as downward nominal wage rigidity or borrowing constraints), since then model-implied impulse responses with respect to positive government spending shocks will differ from those for negative shocks. Supplemental Appendix A.5 gives further examples of weight functions for several identified tax, technology, and monetary policy shocks. As these examples illustrate, plotting estimates of the weights ω_X is useful in interpreting the results of any subsequent impulse response analysis and for comparing with

⁷Pointwise confidence intervals can be obtained with conventional heteroskedasticity-robust standard errors. One could also use autocorrelation robust standard errors to allow for time series dependence of X_t , but causal interpretation is more challenging if the shocks are not independent.

⁸Our code and data are available at https://github.com/mikkelpm/nonlinear_dynamic_causal

EMPIRICAL WEIGHT FUNCTIONS: GOVERNMENT SPENDING SHOCKS




Figure 1: Estimated causal weight functions ω_X for government spending shocks obtained from the replication files for [Ramey \(2016\)](#), quarterly data. Horizontal axis in units of standard deviations. “ $\omega > 0$ ”: total weight $\int_0^\infty \omega_X(x) dx$ on positive shocks. Papers referenced: [Blanchard and Perotti \(2002\)](#), [Fisher and Peters \(2010\)](#), [Ben Zeev and Pappa \(2017\)](#), [Ramey \(2011\)](#).

prior studies.

PARAMETRIC NONLINEAR SPECIFICATIONS. In many cases, economic theory predicts that the average structural function Ψ_h is likely nonlinear. For example, if the outcome variable has limited support, such as due to censoring or when it is discrete, the structural function must necessarily be nonlinear. In such cases, it seems natural to model the nonlinearity directly, rather than to stick to a linear specification as in (7). For example, [Jordà \(2005\)](#) and [Jordà and Taylor \(2025\)](#) suggest including powers of the shock in local projections. Similarly, there is a rich literature on nonlinear extensions of VAR models, see for example [Kilian and Lütkepohl \(2017, Chapter 18\)](#). Such direct modeling of the nonlinearities is sensible if the goal of the analysis is to directly characterize the extent and types of nonlinearity present in the data, e.g., threshold effects or sign and size dependence ([Caravello and Martínez Bruera, 2024](#)).

However, for estimating average causal effects, simple linear local projections or VARs appear more robust than nonlinear parametric specifications. As shown in [Proposition 1](#), the linear specification in (7) is robust to misspecification in that it estimates a well-defined average marginal effect regardless of the form of nonlinearity in the structural function Ψ_h . By contrast, we now show that this is not the case for a local projection specification that includes a quadratic term, echoing similar results in [Angrist \(2001\)](#) regarding robustness of parametric nonlinear limited dependent variable models. These results suggest that nonlinear specifications do not generally have such a robustness property.

Consider a quadratic local projection of Y_{t+h} on X_t , X_t^2 , and an intercept.⁹ Assume for analytical simplicity that X_t has a standard normal distribution, so in particular X_t and X_t^2 are uncorrelated (our qualitative conclusions can be shown to go through without the normality assumption). Then the population version of the projection is

$$Y_{t+h} = \beta_{0,h} + \beta_{1,h}X_t + \beta_{2,h}X_t^2 + \text{residual}_{h,t+h},$$

with implied derivative of the regression function at $X_t = x$ given by

$$\bar{\beta}_h(x) \equiv \beta_{1,h} + 2\beta_{2,h}x, \tag{10}$$

and population regression coefficients

$$\beta_{1,h} \equiv \frac{\text{Cov}(g_h(X_t), X_t)}{\text{Var}(X_t)}, \quad \beta_{2,h} \equiv \frac{\text{Cov}(g_h(X_t), X_t^2)}{\text{Var}(X_t^2)}. \tag{11}$$

Proposition 2. *Assume that $X_t \sim N(0, 1)$, and that g_h defined in (5) is differentiable with a derivative that is locally absolutely continuous on \mathbb{R} . Finally, assume $E[|g_h(X_t)| + |g'_h(X_t)| + |g''_h(X_t)|] < \infty$. Then, using the definitions (10)–(11),*

$$\bar{\beta}_h(x) = E[(1 + X_t x)g'_h(X_t)] = E[g'_h(X_t)] + xE[g''_h(X_t)]. \tag{12}$$

The first expression in (12) shows that the estimated derivative $\bar{\beta}_h(x)$ equals a weighted average of the true derivative function $g'_h(\cdot)$, but with weights that are negative whenever $1 + X_t x < 0$.¹⁰ If the true regression function g_h is in fact quadratic, then $\bar{\beta}_h(x)$ is consis-

⁹While we focus on the quadratic case for simplicity, [Proposition 2](#) below can be shown to generalize to a polynomial specification of any fixed order.

¹⁰It also follows from the proposition that any estimated weighted average derivative $\int \omega(x)\bar{\beta}(x)dx$ that

tent for the marginal effect function $\Psi'_h(x)$. But if the regression function is misspecified, the negative weighting leads to a sign reversal: the second expression in (12) implies that even if g_h is monotonically increasing, the estimated derivative $\bar{\beta}_h(x)$ will be negative for sufficiently large x whenever $E[g''_h(X)] < 0$. Such sign reversal is not shared by the linear estimator (7), for which the weighting scheme ω_X is convex. This lack of robustness of a quadratic (or more generally polynomial) specification of the regression function to functional form misspecification is related to the observation in White (1980) that polynomial approximations to the conditional mean function g_h cannot be interpreted as providing a Taylor series approximation to g_h .

STATE-DEPENDENT SPECIFICATIONS. A particularly popular nonlinear local projection specification in applied work is a state-dependent specification that interacts the shock with a binary regime indicator $S_t \in \{0, 1\}$ (see Cloyne, Jordà, and Taylor, 2023, Gonçalves, Herrera, Kilian, and Pesavento, 2024b, and references therein):

$$Y_{t+h} = \hat{\beta}_{0,h}(1 - S_t)X_t + \hat{\beta}_{1,h}S_tX_t + \hat{\gamma}'_{0,h}(1 - S_t)\mathbf{W}_t + \hat{\gamma}'_{1,h}(S_t\mathbf{W}_t) + \text{residual}_{h,t+h}.$$

For example, S_t may indicate whether the economy is in a recession. Assuming that the local projection is fully interacted as above (i.e., all control variables \mathbf{W}_t are interacted with S_t), then the procedure is tantamount to running separate regressions on the subsamples with $S_t = 0$ and $S_t = 1$, respectively.¹¹ It follows that all the analysis surrounding Proposition 1 above applies upon conditioning on $S_t = s \in \{0, 1\}$. In particular, the probability limit of the state-dependent impulse response estimate $\hat{\beta}_{s,h}$ equals a positively weighted average of conditional marginal effects $\partial E[Y_{t+h} \mid X_t = x, S_t = s] / \partial x$, which have a clear causal interpretation provided the shock independence assumption (2) holds conditional on S_t (i.e., within each regime). Thus, despite their apparent linearity conditional on regime, state-dependent local projections identify causal estimands even when the true DGP has a nonlinear form, such as a model with smooth or discrete regime-switching. However, consistent with the discussion in Section 2, it is important to interpret the impulse responses as averaging over all future shocks, *including* potential future regime switches. In other words, the local projection estimand does not hold the regime fixed within the impulse response horizon.

is a nontrivial function of the coefficient $\beta_{2,h}$ (i.e., whenever $\int x\omega(x)dx \neq 0$) equals a weighted average of $g'_h(\cdot)$ with weights that are negative for some x .

¹¹If we instead omit the interaction terms from the regression and only control linearly for S_t , then we are in the case of Sections 3.3 and 6.2 below.

3.2 Identification with proxies

In many applications, observations of the shock are contaminated by measurement error, such as when accurate measurements are available only in a subset of the time periods. In such cases, researchers typically treat the measurements Z_t as a *proxy* for the shock of interest X_t , or, equivalently, an instrument for the shock (see [Stock and Watson, 2018](#), for a review). We now show that when the structural function is nonlinear, linear VARs and local projections onto the proxy identify average marginal effects up to scale, provided that the conditional mean of the proxy given the shock is monotone in the shock.

We assume that the proxy Z_t is valid, in the sense that it satisfies the exclusion restriction

$$E[Y_{t+h} \mid X_t, Z_t] = E[Y_{t+h} \mid X_t] = g_h(X_t), \quad (13)$$

formalizing the notion that if the shock X_t were observed, the proxy Z_t would not provide any further predictive power. It is implied by the standard assumption in the measurement error literature that the measurement error in Z_t is non-differential, i.e., the whole conditional distribution of Y_{t+h} given (X_t, Z_t) depends only on X_t (or equivalently that Z_t is independent of $\mathbf{U}_{h,t+h}$) (e.g., [Carroll, Ruppert, Stefanski, and Crainiceanu, 2006](#), Chapter 2.6).

We consider the “reduced-form” local projection of the outcome Y_{t+h} on the proxy Z_t . Under (13), the population version of this regression has slope coefficient

$$\tilde{\beta}_h \equiv \frac{\text{Cov}(\zeta(X_t), g_h(X_t))}{\text{Var}(Z_t)}, \quad (14)$$

where

$$\zeta(x) \equiv E[Z_t \mid X_t = x]. \quad (15)$$

As shown by [Plagborg-Møller and Wolf \(2021\)](#), $\tilde{\beta}_h$ also corresponds to the probability limit of an impulse response from a structural VAR where the proxy is ordered first, and the specification controls for sufficiently many lags.

Proposition 3. *Assume that X_t is continuously distributed on an interval $I \subseteq \mathbb{R}$ (the interval may be unbounded, and could equal \mathbb{R}), and that the variance of Z_t is positive and finite. Assume that the conditional mean g_h defined in (5) is locally absolutely continuous on I , and $E[|g_h(X_t)|(1 + |\zeta(X_t)|)] < \infty$. Finally, assume that $\int_I |\tilde{\omega}_Z(x) g'_h(x)| dx < \infty$, where*

$$\tilde{\omega}_Z(x) \equiv \frac{\text{Cov}(\mathbb{1}\{X_t \geq x\}, \zeta(X_t))}{\text{Var}(Z_t)}, \quad (16)$$

and that for sufficiently large positive and negative x , the function $\zeta(x) - E[Z_t]$ does not change sign.¹² Then, the proxy estimand (14) satisfies

$$\tilde{\beta}_h = \int_I \tilde{\omega}_Z(x) g'_h(x) dx.$$

The weight function $\tilde{\omega}_Z$ has the following properties:

- (i) It is equivariant to additive and multiplicative measurement error: If $\tilde{Z}_t = V_{1t} + V_{2t}Z_t$, where (V_{1t}, V_{2t}) is a bivariate random vector independent of (X_t, Z_t) , then $\tilde{\omega}_{\tilde{Z}}(x) = \frac{E[V_{2t}] \text{Var}(Z_t)}{\text{Var}(\tilde{Z}_t)} \tilde{\omega}_Z(x)$ for all x .
- (ii) It is nonnegative, $\tilde{\omega}_Z(x) \geq 0$, provided that $E[Z_t | X_t \geq x] \geq E[Z_t | X_t < x]$.
- (iii) It depends only on the joint distribution of (X_t, Z_t) , but not on the conditional distribution of Y_{t+h} given X_t .

A sufficient condition for property (ii) is that the conditional mean function $\zeta(x)$ is monotone increasing. Under this assumption, $\tilde{\omega}_Z$ is also hump-shaped: monotonically increasing from 0 to its maximum for $x \leq x_0$, and then monotonically decreasing back to 0 for $x \geq x_0$, where $x_0 \equiv \inf\{x \in I : \zeta(x) \geq E[Z_t]\}$.

Combining Proposition 3 with the identification result (6) implies that linear proxy regressions identify weighted averages of marginal effects, $\theta_h(\tilde{\omega}_Z) = \int \tilde{\omega}_Z(x) \Psi'_h(x) dx$, just as in the case of directly observed shocks. Unlike in the observed shocks case, the weights $\tilde{\omega}_Z$ will not be positive unless the proxy satisfies the condition in point (ii) of Proposition 3—this condition is slightly weaker than monotonicity of $\zeta(x)$.¹³ However, monotonicity of ζ ensures not just that the weights are positive, but also that they have an intuitive hump-shape, giving most weight to shocks in the middle of the distribution.

Monotonicity of ζ is implied by, but much weaker than the continuous-treatment version of the Imbens and Angrist (1994) monotonicity condition, needed for causal interpretation of two-stage least squares estimands under endogeneity. We defer the details to Supplemental Appendix A.2, where we generalize the identification results in Angrist, Graddy, and Imbens (2000) by allowing for non-smooth potential outcome functions and non-binary Z_t ; we also

¹²That is, there exist $\underline{x}, \bar{x} \in I$ and $\underline{z}, \bar{z} \in \{-1, 1\}$ such that $\underline{z}\{\zeta(x) - E[Z_t]\} \geq 0$ for all $x \leq \underline{x}$ and $\bar{z}\{\zeta(x) - E[Z_t]\} \geq 0$ for all $x \geq \bar{x}$.

¹³For instance, the condition may still hold even if monotonicity is violated over a sufficiently small interval in the middle of the support of X_t .

extend [Proposition 3](#) to the case with covariates. It follows from this identification result that monotonicity of ζ holds under much weaker conditions than those required for causal interpretation of $\tilde{\beta}_h$ under endogeneity. [Rambachan and Shephard \(2021, Theorem 7\)](#) derive an alternative characterization of the proxy estimand [\(14\)](#) involving derivatives of the *reduced-form* potential outcome as a function of the proxy Z_t (rather than of the shock X_t), and therefore the monotonicity assumption has no counterpart in their analysis.

A practical implication of [Proposition 3](#) is that applied researchers should seek to construct proxies that are credibly positively related to the unobserved latent shock of interest. However, it is not essential that the relationship is linear or indeed of any particular known functional form. Of course, constructing valid proxies may be challenging in practice: measurement error in their construction may fail to be non-differential, leading to violation of the exclusion restriction [\(13\)](#), and lack of variability of Z_t in the data may lead to large standard errors for $\tilde{\beta}_h$. The purpose of [Proposition 3](#) is to clarify the value of this hard work, if done well.

Example 2. An interesting example of a proxy is one constructed from so-called “narrative sign restrictions”, where it is assumed that the researcher observes not the shock itself, but a discrete signal of whether a large shock occurred. While [Antolín-Díaz and Rubio-Ramírez \(2018\)](#) and [Giacomini, Kitagawa, and Read \(2023\)](#) exploit such restrictions in a likelihood framework, [Plagborg-Møller and Wolf \(2021\)](#) and [Plagborg-Møller \(2022\)](#) recommend treating them as a special case of proxy identification.

As a concrete example, assume that for some constants $c_1, c_2 \geq 0$ (which may be unknown to the econometrician), $Z_t = \mathbb{1}\{X_t \geq c_2\} - \mathbb{1}\{X_t \leq -c_1\}$. That is, the proxy equals 1 for sufficiently large positive shocks, -1 for sufficiently large negative shocks, and is otherwise uninformative.¹⁴ Let $F_X(x) \equiv P(X_t \leq x)$ be the cumulative distribution function (CDF) of X_t . Then the weight function $\tilde{w}_Z(x)$ is nonnegative and proportional to

$$\text{Cov}(\mathbb{1}\{X_t \geq x\}, Z_t) = \begin{cases} F_X(x)[2 - F_X(c_2) - F_X(-c_1)] & \text{for } x \leq -c_1, \\ F_X(x)[1 - F_X(c_2) - F_X(-c_1)] + F_X(-c_1) & \text{for } x \in (-c_1, c_2), \\ [1 - F_X(x)][F_X(c_2) + F_X(-c_1)] & \text{for } x \geq c_2, \end{cases}$$

¹⁴This example assumes that we correctly classify *all* episodes with shocks of sufficiently large magnitude. However, [Proposition 3](#) shows that the calculations continue to apply (up to scale) even if there is random misclassification of the form $Z_t = V_t[\mathbb{1}\{X_t \geq c_2\} - \mathbb{1}\{X_t \leq -c_1\}]$, where V_t is a Bernoulli random variable that is independent of (X_t, Y_{t+h}) .

as can be verified through direct calculation. It is easy to see that the above weight function is hump-shaped: monotonically increasing until either $x = -c_1$ or $x = c_2$ (depending on the sign of $1 - F_X(c_2) - F_X(-c_1)$), and then monotonically decreasing. Arguably, such a weight function is economically sensible. In fact, if $1 - F_X(c_2) = F_X(-c_1)$ (as would be the case if $c_1 = c_2$ and the distribution of X_t were symmetric around 0), then the weight function is “nearly” uniform as it is shaped like a plateau: increasing for $x < -c_1$, then flat for $x \in [-c_1, c_2]$, then decreasing.

This example shows that conventional proxy local projections or VARs can estimate meaningful causal summaries even if the proxy (which here is discrete) is quite nonlinearly related to the true (continuous) shock, and in ways that are not directly known to the econometrician. This robustness may not be shared by likelihood-based approaches to identification via narrative restrictions. ■

The weight function (16) does not integrate to 1 due to attenuation bias, so that we only identify average marginal effects up to scale. However, since the weight function doesn’t depend on the outcome, this is not an issue in practice: we can scale $\tilde{\beta}_h$ by the response of some normalization variable to the proxy (this is the so-called unit effect normalization) to identify a *relative* marginal effect. The local projection instrumental variable estimator of Stock and Watson (2018), which is a two-stage least squares version of local projection, automatically performs this normalization.

Since the shock X_t is not directly observed, we cannot generally estimate the weight function $\tilde{\omega}_Z$ in the data. Instead, it may be useful to plot the observed-shock weight function (9) pretending that Z_t is the actual shock of interest. If it happens that $Z_t \approx X_t$, then these weights will be close to the proxy weights $\tilde{\omega}_Z$, so the plot provides a “best-case” scenario.

3.3 Identification with control variables

In applications where it is challenging to isolate purely exogenous shifts in policy or fundamentals, researchers may be willing to assume that the observed variable X_t (which could be a policy instrument) is exogenous conditional on some control variables \mathbf{W}_t (such as variables that comprise the policy-makers information set):

$$X_t \perp\!\!\!\perp \mathbf{U}_{h,t+h} \mid \mathbf{W}_t. \quad (17)$$

This is a selection on observables assumption as in Angrist and Kuersteiner (2011) and Angrist, Jordà, and Kuersteiner (2018). For example, the assumption holds if $X_t = \Upsilon(\varepsilon_t, \mathbf{W}_t)$,

where ε_t is a shock that is independent of $(\mathbf{W}'_t, \mathbf{U}'_{h,t+h})'$, a nonparametric version of the recursive (or Cholesky) assumption in linear structural VAR identification (e.g., [Christiano, Eichenbaum, and Evans, 1999](#)). Then the conditional expectation function

$$g_h(x, \mathbf{w}) \equiv E[Y_{t+h} \mid X_t = x, \mathbf{W}_t = \mathbf{w}]$$

equals the *conditional* average structural function in the causal model (1):

$$g_h(x, \mathbf{w}) = E[\psi_h(x, \mathbf{U}_{h,t+h}) \mid \mathbf{W}_t = \mathbf{w}] \equiv \Psi_h(x, \mathbf{w}),$$

where the expectation is taken with respect to the conditional distribution of the nuisance shocks $\mathbf{U}_{h,t+h}$ given \mathbf{W}_t .

Even under the selection on observables assumption (17), the local projection with controls (7) need not estimate an average marginal effect if the relationship between X_t and the controls is nonlinear. This result extends to recursively identified structural VARs, due to the nonparametric equivalence between these procedures ([Plagborg-Møller and Wolf, 2021](#)). [Proposition 7](#) below shows that the population local projection coefficient β_h can still be written as a weighted average of the marginal effects $\partial g(x, \mathbf{w})/\partial x = \partial \Psi_h(x, \mathbf{w})/\partial x$, but the weights can be negative if the true “propensity score” $\pi^*(\mathbf{w}) \equiv E[X_t \mid \mathbf{W}_t = \mathbf{w}]$ is nonlinear (or, in other words, if the purported “shock” that is the residual from a linear projection X_t on the controls \mathbf{W}_t is predictable by *nonlinear* transformations of the controls). We leave the details, which extend the analysis of [Goldsmith-Pinkham, Hull, and Kolesár \(2024\)](#) to cases with non-discrete X_t , to [Section 6](#). As usual, negative weights are worrying, as they may lead to a sign-reversal. If \mathbf{W}_t just consists of a set of mutually exclusive dummies, then linearity of the propensity score comes for free, and the weights are guaranteed to be positive. However, outside this case, we recommend that researchers do careful sensitivity checks with respect to both the set of controls and the functional form for the controls (e.g., whether the addition of nonlinear transformations of the controls, such as interactions and polynomials, improves prediction of X_t). More research is warranted on best practices for such sensitivity checks, given the data limitations in applied macroeconomics.

4 The bad: identification via heteroskedasticity

Identification via heteroskedasticity has become a popular procedure for causal identification in applications where direct shock measures are unavailable, following [Sentana and Fioren-](#)

tini (2001), Rigobon (2003), and Rigobon and Sack (2004).¹⁵ In a pair of highly-cited papers, Lewbel (2012, 2018) exploits this idea to achieve identification in cross-sectional regressions with endogenous variables and no external instruments (see also Klein and Vella, 2010, for a related approach). In stark contrast to Section 3, the results in this section deliver *bad* news regarding the sensitivity of identification approaches via heteroskedasticity to the assumption that the underlying structural function is linear: the Rigobon-Sack-Lewbel estimator does not generally estimate average marginal effects; more generally, we show that the nonparametric analogue of the identification approach yields very large identified sets for causal effects. However, on the positive side, it is possible to test the linearity assumption in the data.

4.1 Nonparametric version of the identification approach

To explain the sensitivity of conventional identification via heteroskedasticity to the linearity assumption on the structural function, it is helpful to first lay out a nonparametric version of the framework before we review the linear case. Since this section is mainly concerned with giving examples of how the identification approach can fail, we specialize the dynamic set-up from Section 2 to a simpler static model.

NONPARAMETRIC SETUP. We observe an n -dimensional vector \mathbf{Y} of variables that are nonlinearly related to a latent, scalar shock of interest X as well as an $(m - 1)$ -dimensional latent vector \mathbf{U} of nuisance shocks:

$$\mathbf{Y} = \boldsymbol{\psi}(X, \mathbf{U}), \quad X \perp\!\!\!\perp \mathbf{U}, \quad (18)$$

where we suppress time subscripts to ease notation. The above model is a (static) *nonparametric factor model*, since we do not impose parametric restrictions on the unknown structural function $\boldsymbol{\psi}: \mathbb{R}^m \rightarrow \mathbb{R}^n$.

The econometrician observes a scalar D that is informative about the heteroskedasticity of the shock of interest X but independent of the nuisance shocks \mathbf{U} (jointly with X):

$$(D, X) \perp\!\!\!\perp \mathbf{U}. \quad (19)$$

¹⁵Similar identification approaches were developed in the signal-processing literature in the 1990s, see the review by Hyvärinen, Karhunen, and Oja (2001, Section 18.2).

This assumption implies that D is a valid proxy for X , in the sense that \mathbf{Y} and D are independent conditional on X . But because the variable D only influences the variance and higher moments of X but not its mean,

$$E[X \mid D] = 0, \quad (20)$$

we cannot use the proxy in local projections as in [Section 3.2](#). For concreteness, it may be useful to think of D as a binary *regime indicator*, which affects the conditional variance $\text{Var}(X \mid D)$ but not the conditional mean (20), as in the original work by [Rigobon \(2003\)](#).

If we assume that the structural function ψ is linear, it is possible to achieve identification even if D is unobserved, and we relax (19) by allowing D to affect the variances of nuisance shocks. See [Bacchiocchi, Bastianin, Kitagawa, and Mirto \(2024\)](#) and [Lewis \(2024, Section 3\)](#) for excellent reviews. However, since we are only interested in showing how the basic identification approach can fail in a nonparametric context, we maintain the stronger assumptions above. It then follows *a fortiori* that nonparametric identification is even more challenging under weaker assumptions.

REVIEW OF LINEAR IDENTIFICATION. If the structural function ψ in (18) is known to be partially linear, identification of causal effects obtains under an additional relevance assumption. Thus, we temporarily assume that

$$\psi(x, \mathbf{u}) = \boldsymbol{\theta}x + \boldsymbol{\gamma}(\mathbf{u}), \quad (21)$$

where $\boldsymbol{\theta}$ is the unknown vector of causal effects of X , while $\boldsymbol{\gamma}: \mathbb{R}^{m-1} \rightarrow \mathbb{R}^n$ is an unknown function. Following [Rigobon and Sack \(2004\)](#) and [Lewbel \(2012\)](#), construct the scalar instrumental variable

$$Z \equiv (D - E[D])Y_1, \quad (22)$$

where Y_1 is the first element of \mathbf{Y} . In applications, Y_1 may be a policy instrument that is known to be strongly related to X , though it is also allowed to be correlated with the nuisance shocks. Under the linear model (21) and the identification assumptions (19)–(20), Z satisfies the exogeneity restriction for linear identification in [Stock and Watson \(2018\)](#) since $E[Z \mid \mathbf{U}] = 0$. In particular, under these assumptions, a regression of \mathbf{Y} on Y_1 using Z

as an instrument identifies the (relative) causal effects of X :

$$\frac{1}{\text{Cov}(Y_1, Z)} \text{Cov}(\mathbf{Y}, Z) = \frac{1}{\theta_1} \boldsymbol{\theta}. \quad (23)$$

To ensure we are not dividing by zero, we need to additionally assume the relevance conditions that (i) the shock of interest is heteroskedastic across regimes, $\text{Cov}(X^2, D) \neq 0$, and (ii) the causal effect of X on Y_1 is nonzero, $\theta_1 \neq 0$. For completeness, we review the calculations leading to (23) in Supplemental Appendix A.3.

4.2 Fragility under nonlinearity

We now argue that the simple linear identification argument fundamentally cannot be extended to nonparametric contexts.

NONPARAMETRIC IDENTIFIED SET. We first show that the nonparametric model of identification via heteroskedasticity yields a large identified set for the causal effects of X on \mathbf{Y} . To do this, we strengthen the independence and conditional mean assumptions (19)–(20) by imposing a specific model for the relationship between X and D :

$$X = \sigma(D)R, \quad \text{where } R, D, \text{ and } \mathbf{U} \text{ are mutually independent,} \quad (24)$$

$\sigma: \mathbb{R} \rightarrow \mathbb{R}_+$ is a *known* function, and R has a *known* distribution that is symmetric around 0. This model would, for example, be consistent with the conditionally Gaussian model $X \mid D \sim N(0, \sigma^2(D))$.

Proposition 4. *Assume that $(\mathbf{Y}, D, R, X, \mathbf{U})$ satisfy the nonparametric factor model (18) and identification assumption (24). Then there exists an alternative structural function $\tilde{\boldsymbol{\psi}}: \mathbb{R}^2 \rightarrow \mathbb{R}^n$ and a scalar random variable \tilde{U} independent of (R, D, X) such that $(\tilde{\mathbf{Y}}, D)$ has the same joint distribution as (\mathbf{Y}, D) , where*

$$\tilde{\mathbf{Y}} \equiv \tilde{\boldsymbol{\psi}}(X, \tilde{U}),$$

and such that $\tilde{\boldsymbol{\psi}}(-x, \tilde{u}) = \tilde{\boldsymbol{\psi}}(x, \tilde{u})$ for all x, \tilde{u} .

The proposition states that the identified set for $\boldsymbol{\psi}$ is so large that it always contains a structural function $\boldsymbol{\psi}(x, \mathbf{u})$ that is symmetric in x around 0. In particular, we can never rule out that the average marginal effect $\int \omega(x)(\partial E[\boldsymbol{\psi}(x, \mathbf{U})]/\partial x) dx$ is zero when the weight

function $\omega(x)$ is symmetric around 0. Intuitively, the challenge is that D does not affect the mean of X , only higher moments, so—without strong functional form restrictions on the relationship between the outcomes and the shocks—we do not have enough information to sign mean effects of shifts in the latent shock X . This holds even though we assume that the econometrician knows exactly how D affects the dispersion of the X distribution. Notice that the construction of the observationally equivalent symmetric structural function in [Proposition 4](#) only relies on a single (scalar) nuisance shock; hence, knowledge about the true number of shocks does not ameliorate the identification failure (see [Section 5](#) for further discussion of this point).

A careful inspection of the proof of [Proposition 4](#) reveals that the result is closely related to a known issue with identification via heteroskedasticity in a linear context: while the variance of the shock of interest X must vary across regimes, we cannot simultaneously allow the impulse responses of X to vary across regimes (see [Lewis, 2024](#), Section 6.1, for a discussion and references). However, in a nonparametric context this problem is even worse, since there is no fundamental distinction between “coefficients” and “shock variances” in a general non-linear model. *A priori* restrictions that certain “coefficients” are independent of the regime are only meaningful once we parametrize the model, which complicates the development of an empirically useful nonparametric generalization of the identification approach.

SENSITIVITY OF LINEAR PROCEDURES. Because the nonparametric identified set is large, we can expect estimation procedures based on linearity of the structural function to fail to estimate causal objects in general. The next result implies that this is indeed the case for the linear instrumental variable estimator (23) of [Rigobon and Sack \(2004\)](#) and [Lewbel \(2012\)](#).

Proposition 5. *Assume the additively separable structural model*

$$\mathbf{Y} = \boldsymbol{\Psi}(X) + \boldsymbol{\gamma}(\mathbf{U}),$$

where $\boldsymbol{\Psi}: \mathbb{R} \rightarrow \mathbb{R}^n$, $\boldsymbol{\gamma}: \mathbb{R}^{m-1} \rightarrow \mathbb{R}^n$, and we normalize $E[\boldsymbol{\Psi}(X)] = E[\boldsymbol{\gamma}(\mathbf{U})] = \mathbf{0}$. Suppose that the independence assumption (19) holds, and let Z be given by (22). Suppose that the variables (\mathbf{Y}, Z, D) have finite second moments, and that the support of X is given by the interval $I \subseteq \mathbb{R}$ (the interval may be unbounded, and could equal \mathbb{R}). Suppose also that for each j , Ψ_j (the j -th component of $\boldsymbol{\Psi}$) is locally absolutely continuous on I , and that for some

$\underline{x}, \bar{x} \in I$, $\Psi_j(x)$ is monotone for $x \leq \underline{x}$ and for $x \geq \bar{x}$. Then

$$\text{Cov}(\mathbf{Y}, Z) = \int \check{\omega}(x) \Psi'(x) dx,$$

where

$$\check{\omega}(x) \equiv \text{Cov} \left(\mathbb{1}\{X \geq x\}, \Psi_1(X)(D - E[D]) \right). \quad (25)$$

Proposition 5 shows that regressing \mathbf{Y} onto the instrument Z yields a weighted average of marginal effects, but with a weight function $\check{\omega}(x)$ that cannot be guaranteed to be positive.¹⁶ In fact, the weights even integrate to 0 in some cases, for example if $\Psi_1(x) = \Psi_1(-x)$ and the conditional distribution of X given D is symmetric around 0. In such cases, the instrument erroneously estimates a zero causal effect of X on Y_j for $j \geq 2$ even if $\Psi_j(x) = \theta_j x$ is a linear function with $\theta_j \neq 0$.

The weights can also be negative—and therefore cause the econometrician to get the sign of the marginal effects wrong—even in the seemingly favorable setting where (unbeknownst to the econometrician) the policy variable Y_1 simply equals the shock of interest X , without any nonlinearity or contamination by nuisance shocks, i.e., $\Psi_1(x) = x$. Assume in addition, as in [Rigobon \(2003\)](#), that the regime indicator $D \in \{0, 1\}$ is binary and $E[X | D] = 0$. Then a simple calculation shows that the weights in [equation \(25\)](#) equal

$$\check{\omega}(x) = \text{Var}(D) \int_x^\infty [f_{X|1}(v) - f_{X|0}(v)] v dv = \text{Var}(D) \int_{-\infty}^x [f_{X|0}(v) - f_{X|1}(v)] v dv, \quad (26)$$

where $f_{X|d}(x)$ is the density of X conditional on regime $D = d$. Suppose the right (resp., left) tail of the X distribution is fatter (resp., thinner) in regime $D = 1$ than in regime $D = 0$, meaning that $f_{X|1}(x) > f_{X|0}(x)$ for $x \gg 0$ and $f_{X|0}(x) > f_{X|1}(x)$ for $x \ll 0$. Then it follows from [equation \(26\)](#) that $\check{\omega}(x) > 0$ for $x \gg 0$, while $\check{\omega}(x) < 0$ for $x \ll 0$. This simple example shows that the instrumental variable estimator can easily generate negative weights, even when it satisfies the exclusion and relevance conditions and the policy variable is linear in the shock. To trust that the weights are positive, we would need to have quite detailed information about the conditional shock density in the two regimes; simple moment restrictions do not suffice.

Intuitively, the problem of negative weights comes about because the [Rigobon \(2003\)](#) and [Lewbel \(2012\)](#) instrumental variable Z defined in [\(22\)](#) fails the proxy monotonicity

¹⁶This is not a special case of [Proposition 3](#), since Z does not satisfy the nonparametric proxy assumption [\(13\)](#).

assumption discussed earlier in connection with [Proposition 3](#). Because the only source of exogenous variation is the regime indicator D , and this indicator does not affect the mean of the latent shock X but only higher moments, it is generally impossible to construct any proxy variable that is guaranteed to be monotone in X , unless we make strong assumptions about the structural function.

If the model is not additively separable as assumed in [Proposition 5](#), the instrumental variables estimator can exhibit even more pathological behavior, in that it may not equal a weighted average of marginal effects at all. As a simple example, consider a multiplicative model $\mathbf{Y} = X\boldsymbol{\gamma}(\mathbf{U})$ with $E[\boldsymbol{\gamma}(\mathbf{U})] = \mathbf{0}$ and impose the independence assumption [\(19\)](#). In that model, $E[\mathbf{Y} \mid X] = \mathbf{0}$, so the marginal effect function is identically zero, but $\text{Cov}(\mathbf{Y}, Z) = \text{Cov}(X^2, D) \text{Cov}(\boldsymbol{\gamma}_1(\mathbf{U}), \boldsymbol{\gamma}(\mathbf{U})) \neq 0$ in general, so the instrument erroneously estimates a nonzero effect.

4.3 Silver lining: testability of the linearity assumption

While the sensitivity of identification via heteroskedasticity to linearity of the structural function $\boldsymbol{\psi}$ is disheartening, at least the linear model [\(21\)](#) implies testable restrictions. Specifically, as noted by [Rigobon and Sack \(2004\)](#) and [Wright \(2012\)](#), for any d_0, d_1 in the support of D , the difference $\text{Var}(\mathbf{Y} \mid D = d_1) - \text{Var}(\mathbf{Y} \mid D = d_0) = [\text{Var}(X \mid D = d_1) - \text{Var}(X \mid D = d_0)]\boldsymbol{\theta}\boldsymbol{\theta}'$ should be a rank-1 matrix under linearity and the maintained independence assumption [\(19\)](#). Other over-identification tests in more general linear models of identification via heteroskedasticity are discussed in the review article by [Lewis \(2024\)](#). We are not aware of any thorough analysis of the power properties of these tests against nonlinear alternatives.

5 The ugly: identification via non-Gaussianity

A second approach to identification in linear models in the absence of direct shock measures is to assume that the structural shocks are mutually independent and non-Gaussian; see [Gouriéroux, Monfort, and Renne \(2017\)](#), [Lanne, Meitz, and Saikkonen \(2017\)](#), and the review article by [Lewis \(2024\)](#). [Lewbel, Schennach, and Zhang \(2024\)](#) propose a similar approach to achieve identification in cross-sectional endogenous regressions in the absence of external instruments. An earlier literature outside economics goes by the name independent components analysis (ICA), see [Kagan, Linnik, and Rao \(1973, Chapter 10\)](#), [Comon \(1994\)](#), and the textbook by [Hyvärinen, Karhunen, and Oja \(2001\)](#).

This section delivers *ugly* news regarding the sensitivity of this identification approach to linearity of the structural function: once the linearity assumption is dropped, the non-Gaussianity assumption is essentially vacuous; as a consequence, estimators based on non-Gaussianity and linearity of the structural function can fail spectacularly even under mild departures from linearity. What is worse, the linearity assumption is untestable in general.

5.1 Nonparametric version of the identification approach

As in [Section 4](#), we consider the nonparametric factor model [\(18\)](#). However, unlike in the case of identification via heteroskedasticity, we now do not observe any additional proxy variables that aid in identifying the latent shocks. Instead, we hope to achieve identification via restrictions on the distributions of the shocks.

REVIEW OF LINEAR IDENTIFICATION. Assume temporarily that the number of shocks equals the number of observables, $m = n$, and that the structural function is linear:

$$\psi(x, \mathbf{u}) = \boldsymbol{\theta}x + \boldsymbol{\gamma}\mathbf{u}, \quad \boldsymbol{\theta} \in \mathbb{R}^n, \boldsymbol{\gamma} \in \mathbb{R}^{n \times (n-1)}.$$

Assume also that the n shocks (X, U_1, \dots, U_{n-1}) are mutually independent, and at most one of these shocks has a Gaussian distribution. We will refer to this model as the linear ICA model. A deep result in probability theory, the Darmois-Skitovich Theorem, says that two nontrivial linear combinations of independent variables cannot themselves be independent, unless all the underlying variables are Gaussian. In the context of the linear ICA model, the theorem implies that any two linear combinations $\boldsymbol{\varsigma}'\mathbf{Y}$ and $\tilde{\boldsymbol{\varsigma}}'\mathbf{Y}$ of the data $\mathbf{Y} = \boldsymbol{\theta}X + \boldsymbol{\gamma}\mathbf{U}$ can be independent if and only if these linear combinations equal two different shocks in the model (up to sign and scale). Hence, the shocks in the model can be identified by searching for those linear combinations of the observed variables that are independent; once we have the shocks, we can then estimate their causal effects. See [Hyvärinen, Karhunen, and Oja \(2001\)](#) and [Lewis \(2024, Section 4\)](#) for reviews of estimation procedures.

The abstract identification argument above can be made less mysterious through a method of moments framework that exploits implications of shock independence for higher moments of the data ([Lewis, 2024, Section 4.4](#)). Nevertheless, it is clear from both the abstract argument and the more concrete moment-based approach that linearity of the structural function is being leveraged heavily.

5.2 Fragility under nonlinearity

NONPARAMETRIC IDENTIFIED SET. Unfortunately, the mere assumptions that the latent shocks are independent and non-Gaussian provide essentially no identification power in a nonparametric context. The identified set under these assumptions is so large that nearly any function of the data can be labeled a “shock”.

Proposition 6. *Let $\tilde{\mathbf{Y}} = \Upsilon(\mathbf{Y})$ be a homeomorphic¹⁷ transformation of \mathbf{Y} , with j -th element denoted by \tilde{Y}_j . For all $j = 2, \dots, n$, assume that the quantile function of \tilde{Y}_j conditional on $\tilde{Y}_{j-1}, \tilde{Y}_{j-2}, \dots, \tilde{Y}_1$ is continuous in the quantile and the conditioning arguments. Define $\tilde{X} \equiv \tilde{Y}_1$, and let $\{\bar{U}_j\}_{j=1}^{n-1}$ be mutually independent uniform variables on $[0, 1]$ that are also independent of \tilde{X} . Then there exists a continuous function $\bar{\psi}: \mathbb{R}^n \rightarrow \mathbb{R}^n$ such that the random vector*

$$\bar{\mathbf{Y}} \equiv \bar{\psi}(\tilde{X}, \bar{U}_1, \dots, \bar{U}_{n-1})$$

has the same distribution as \mathbf{Y} .

Proposition 6 shows that the nonparametric factor model (18) is very under-identified, even if we restrict the number of latent, independent shocks to equal $m = n$ and impose smoothness on the structural function ψ . Indeed, any element of almost any one-to-one transformation $\Upsilon(\mathbf{Y})$ of the observables could be construed as a “shock” for *some* data-consistent choice of structural function ψ . Note that the identification issues go well beyond the familiar “labeling problem” in linear ICA analysis, where the shocks are identified up to sign and permutation, such that additional economic information is required to label each of the statistically identified shocks (Lewis, 2024, Section 6.4). While the challenge of identifying nonlinear factor models is well known in the broader literature (see the review by Jutten and Karhunen, 2003), it appears that the serious consequences of this fact for shock identification in macroeconometrics have not been explored previously.

The fundamental issue is that non-Gaussianity of the shocks is a vacuous assumption in the nonparametric setting: it is an innocuous normalization to assume that all shocks have uniform distributions, since we can always nonlinearly transform any shock distribution to the uniform distribution via the quantile function. In other words, we have severe identification failure as in Proposition 6 even if the econometrician knows the exact distributions of each shock. Hence, it is no accident that the identification argument in the ICA and

¹⁷That is, continuous, one-to-one, and with a continuous inverse function $\Upsilon^{-1}(\cdot)$.

structural VAR literatures relies heavily on the linearity assumption: there is no nonlinear equivalent of the Darmon-Skitovich Theorem.

If we allow for slightly less smoothness of the structural function ψ , then the identification problem is even worse. As [Gunsilius and Schennach \(2023\)](#) note, any n -dimensional vector \mathbf{Y} can be represented as a nonlinear factor model (18) in a *single* latent shock X (so $m = 1$ and $\mathbf{U} = \mathbf{0}$) using a so-called space-filling curve (e.g., Hilbert curve) construction, though the associated ψ function would not be one-to-one. Hence, without restrictions on the structural function, we cannot rule out that the latent shock of interest X drives *all* the variation in the n observed variables \mathbf{Y} . Indeed, given a uniform random variable U on $[0, 1]$, we can generate an infinite number of independent uniform random variables $\{V_j\}$ from the decimal expansion of $U = 0.u_1^1 u_2^1 u_1^2 u_3^1 u_2^2 u_1^3 u_4^1 u_3^2 u_2^3 u_1^4 \dots$, and taking $V_j \equiv 0.u_1^j u_2^j \dots$ (and hence an infinite number of independent random variables with arbitrary distributions F_j by taking the inverse transform $F_j^{-1}(V_j)$).¹⁸

SENSITIVITY OF LINEAR PROCEDURES. Due to the nonparametric identification failure, we can expect identification approaches based on non-Gaussianity to be very sensitive to exact linearity in the structural function. The following two simple examples illustrate such sensitivity in two settings where the linearity assumption is untestable. Thus, identification via non-Gaussianity is not only fragile, it is also generally not falsifiable.

Example 3. Let the two latent shocks (X, U) have a bivariate standard normal distribution. Let the two observed variables be given by

$$Y_1 \equiv X + U, \quad Y_2 \equiv \gamma(X - U),$$

for an arbitrary measurable nonlinear function $\gamma: \mathbb{R} \rightarrow \mathbb{R}$. We can interpret this setting as being *almost* a linear ICA model, except that the second variable has not been transformed quite correctly. In the above model, Y_1 and Y_2 are independent, and Y_2 has a non-Gaussian distribution.¹⁹ Hence, any linearity-based ICA procedure applied to the data (Y_1, Y_2) will erroneously conclude that the first variable equals the first shock and the second variable the second shock (up to the mean). Moreover, there is nothing in the data that can reject the validity of the linear ICA assumptions. Notice the lack of continuity: even if $\gamma(\cdot)$ is only

¹⁸This is a consequence of the fact that there exists a one-to-one function ϕ such that both ϕ and ϕ^{-1} are measurable between the measurable spaces (M, \mathcal{B}_M) and $([0, 1], \mathcal{B}_{[0,1]})$, where M is any separable complete metric space and \mathcal{B}_M is the Borel σ -algebra on M ([Dudley, 2002](#), Theorem 3.1.1).

¹⁹This is because the vector $(X + U, X - U)$ has a joint normal distribution with uncorrelated components.

slightly nonlinear, linear ICA procedures will conclude (asymptotically) that the first shock contributes 100% of the variance of Y_1 , even though the true number is 50%.

This example illustrates how getting the transformation of Y_2 slightly wrong can mess up causal inference about the other variable Y_1 (which is in fact linear in the true shocks). ■

Example 4. Consider a model of the form

$$Y_1 = X + U, \quad Y_2 = X + \gamma(U),$$

where X and U are independent latent shocks. Supplemental Appendix A.4 gives concrete choices of non-Gaussian distributions of the shocks and a smooth $\gamma(\cdot)$ function such that Y_1 and Y_2 are independent and both non-normal. Hence, as in the previous example, any linearity-based ICA procedure applied to the data (Y_1, Y_2) will erroneously attribute all variation in Y_1 to the first shock and all variation in Y_2 to the second shock. Note that in this example, both of the true shocks (X, U) are non-Gaussian, and the only nonlinearity in the true structural model is the relationship between Y_2 and U . ■

In conclusion, linearity-based ICA identification procedures can be highly misleading under departures from a linear model, as with identification via heteroskedasticity (but unlike identification with observed shocks or proxies). In fact, the situation is arguably worse than in [Section 4](#), since even arbitrarily small structural nonlinearities can yield large biases, and the linearity assumption is not testable in general.

6 Identification of average marginal effects

In [Section 3](#), we reverse-engineered the weight function allowing us to interpret the linear local projection estimand as an average marginal effect. We now consider a *forward-engineering* problem: how can we estimate an average marginal effect with a *pre-specified* weight function? For simplicity, we focus on the case when the shock of interest is observed, but we drop the requirement that the shock is continuously distributed. We first consider the case without control variables before extending the analysis to allow for controls.

6.1 Identification without controls

We consider the setup from [Section 3.1](#), but drop time subscripts to make it clearer that our analysis applies to cross-sectional as well as time series settings. Indeed, part of our goal is

to show that identification arguments from cross-section settings carry directly through to time-series contexts. Let

$$g(x) \equiv E[Y \mid X = x] \quad (27)$$

denote the conditional mean function from a nonparametric regression of the scalar outcome Y onto the scalar variable X . We do not restrict the marginal distribution of X : it can be continuous, discrete, or mixed. Let $I \subseteq \mathbb{R}$ denote a (possibly unbounded) interval that contains the support of X . We are interested in summarizing g by reporting its weighted average derivative, weighted by some pre-specified weight function ω . With some abuse of notation, we still denote this weighted average derivative by

$$\theta(\omega) \equiv \int_I \omega(x) g'(x) dx,$$

as in [Section 2](#), even though we don't require that $g(x)$ corresponds to some structural function. To ensure that this object is well-defined, we assume that g is locally absolutely continuous on I . Since (27) only defines g on the support of X , this requires us to extend g to all of I in cases when there are gaps in the support of X , such as when X is discrete. This can be done by linear interpolation: if $P(X \in (a, b)) = 0$ for some $(a, b) \subseteq I$, we set $g(x) = (g(b) - g(a))(x - a)/(b - a) + g(a)$ for $x \in (a, b)$. If the distribution of X is discrete, this defines the derivative g' as the slope between adjacent support points (and the extension will automatically be locally absolutely continuous provided that the spacing between adjacent support points is bounded away from 0).

A regression-based approach to estimating $\theta(\omega)$ first estimates the entire derivative function $g'(\cdot)$ nonparametrically (by, say, series or kernel regression), and then averages it using the weights ω . The next result shows that we can alternatively estimate $\theta(\omega)$ as a weighted average of outcomes, $\theta(\omega) = E[\alpha(X)Y]$, where α is the Riesz representer of the linear functional $g \mapsto \theta(\omega)$.

Lemma 1. *Let $\omega(x) \equiv E[\mathbb{1}\{X \geq x\}\alpha(X)]$. Suppose that (i) the support of X is contained in a (possibly unbounded) interval $I \subseteq \mathbb{R}$; (ii) g is locally absolutely continuous on I ; (iii) $E[|\alpha(X)|(1 + |g(X)|)] < \infty$ with $E[\alpha(X)] = 0$; and (iv) there exists $x_0 \in I$ such that $E[|\alpha(X) \int_{x_0}^X |g'(x)| dx|] < \infty$. Then*

$$E[\alpha(X)g(X)] = \int_I \omega(x) g'(x) dx. \quad (28)$$

Analogous representations for $\theta(\omega)$ are well-known in the literature if we additionally

assume that X is continuously distributed (e.g., [Newey and Stoker, 1993](#), equation 2.6). The representation is usually derived by directly applying integration by parts. Our proof instead generalizes the proof of Stein’s lemma ([Stein, 1981](#), Lemma 1), allowing us to drop the requirement that X is continuously distributed and impose only very mild regularity conditions, which essentially just require that both sides of [equation \(28\)](#) are well-defined. In particular, absolute continuity of g is needed to ensure that $\theta(\omega)$ is well-defined, and in [Lemma B.1](#) in [Supplemental Appendix B](#), we show that if the tails of ω or the tails of g are monotone, then condition (iv) of [Lemma 1](#) holds provided the integral on the right-hand side of [\(28\)](#) exists.

[Lemma 1](#) gives a recipe for constructing weighting-based estimators of $\theta(\omega)$ for particular choices of weight function ω by replacing the expectation in [equation \(28\)](#) with a sample average and, if the function α is unknown, replacing α with an estimate.²⁰ For instance, suppose X is continuous, and let $\omega(x) = f_X(x)$ correspond to the density of X , so that $\theta(\omega) = E[g'(X)]$ is the (unweighted) average derivative. Then the required weighting is given by $\alpha(x) = -f'_X(x)/f_X(x)$, leading to the estimator of [Härdle and Stoker \(1989\)](#) if one uses kernel estimators to estimate the density and its derivative. If the identification condition [\(6\)](#) holds, this estimator identifies the average causal impact of increasing X by an infinitesimal amount. To estimate the average impact of increasing X by a fixed amount δ (i.e., the unweighted average causal effect), which corresponds to setting $\omega(x) = \frac{1}{\delta} \int_{x-\delta}^x f_X(x) dx$, let $\alpha(x) = -\frac{f_X(x) - f_X(x-\delta)}{\delta f_X(x)}$, replacing the derivative of the density by a discrete change. If we set $\omega(x) = f_X^2(x)$, so that $\alpha(x) = -2f'_X(x)$, and we use a leave-one-out kernel estimator for the derivative of the density, we recover the famous density-weighted average derivative estimator of [Powell, Stock, and Stoker \(1989\)](#). Finally, for $\omega(x) = E[\mathbb{1}\{X \geq x\}X]$, we get $\alpha(x) = x$, and [Lemma 1](#) reduces to [Proposition 1](#): this weighting scheme can be estimated by linear regression. As the proofs of [Propositions 2, 3](#) and [5](#) reveal, these results are also special cases of [Lemma 1](#).²¹

²⁰As we discuss in [Section 6.2](#) below, recent results in the semiparametric literature suggest that rather than picking between this weighting-based approach and the regression-based approach to estimation of $\theta(\omega)$, it may pay off to combine them, yielding a “doubly-robust” estimator.

²¹As a consequence, the assumption that X_t be continuous in [Propositions 1](#) and [3](#) can be dropped.

6.2 Identification with control variables

We now generalize the setup to allow for a vector of controls \mathbf{W} . Consider the weighted average derivative

$$\theta(\omega) \equiv E \left[\int_{I_{\mathbf{W}}} \omega(x, \mathbf{W}) g'(x, \mathbf{W}) dx \right],$$

where the expectation is over the marginal distribution of \mathbf{W} , and g' is the derivative with respect to x of the conditional mean function $g(x, \mathbf{w}) \equiv E[Y \mid X = x, \mathbf{W} = \mathbf{w}]$.²² To ensure this object is well-defined, we assume that for each \mathbf{w} , the weights ω are zero outside the interval $I_{\mathbf{w}}$ containing the conditional support of X given $\mathbf{W} = \mathbf{w}$, and that we can extend $g(\cdot, \mathbf{w})$ to $I_{\mathbf{w}}$ such that $g(\cdot, \mathbf{w})$ is locally absolutely continuous on $I_{\mathbf{w}}$, such as by linearly interpolating across any gaps in the support.

Like in the case without covariates, a regression-based estimator of $\theta(\omega)$ first estimates the derivative of the regression function $g'(x, \mathbf{w})$, and then averages the estimated derivative function using the weights ω and the marginal distribution of the covariates. The next result shows that we can alternatively estimate $\theta(\omega)$ by taking weighted averages of the outcome.

Lemma 2. *Let $\omega(x, \mathbf{w}) \equiv E[\mathbb{1}\{X \geq x\} \alpha(X, \mathbf{W}) \mid \mathbf{W} = \mathbf{w}]$. Suppose that conditional on \mathbf{W} , the following holds almost surely: (i) the support of X is contained in a (possibly unbounded) interval $I_{\mathbf{W}} \subseteq \mathbb{R}$; (ii) $g(\cdot, \mathbf{W})$ is locally absolutely continuous on $I_{\mathbf{W}}$; and (iii) $E[\alpha(X, \mathbf{W}) \mid \mathbf{W}] = 0$. Suppose also that (iv) there exists a function $x_0(\mathbf{W}) \in I_{\mathbf{W}}$ such that $E[|\alpha(X, \mathbf{W}) \int_{x_0(\mathbf{W})}^X |g'(x, \mathbf{W})| dx|] < \infty$; and that (v) $E[|\alpha(X, \mathbf{W})|(1 + |g(X, \mathbf{W})|)] < \infty$. Then*

$$E \left[\int_{I_{\mathbf{W}}} \omega(x, \mathbf{W}) g'(x, \mathbf{W}) dx \right] = E [\alpha(X, \mathbf{W}) g(X, \mathbf{W})]. \quad (29)$$

As discussed in [Section 6.1](#), the representation (29) is well-known if the distribution of X is continuous conditional on \mathbf{W} . The novelty of [Lemma 2](#) is to drop the continuity requirement and relax the regularity conditions.

If $X \in \{0, 1\}$ is a binary treatment variable, and we additionally assume that X is as good as randomly assigned conditional on \mathbf{W} , then $g(1, \mathbf{w}) - g(0, \mathbf{w})$ corresponds to the conditional average treatment effect (ATE) for individuals with $\mathbf{W} = \mathbf{w}$. In this case, the average derivative simplifies to

$$\theta(\omega) = E \left[(g(1, \mathbf{W}) - g(0, \mathbf{W})) \int_{I_{\mathbf{W}}} \omega(x, \mathbf{W}) dx \right]$$

²²Weighting by the marginal distribution of \mathbf{W} is not restrictive, since weighting schemes that use other forms of averaging across \mathbf{w} can be recovered by defining ω appropriately.

$$= E[(g(1, \mathbf{W}) - g(0, \mathbf{W}))\alpha(1, \mathbf{W})P(X = 1 | \mathbf{W})],$$

which corresponds to a weighted average of conditional ATEs. By letting $\alpha(X, \mathbf{W}) = X/P(X = 1 | \mathbf{W}) - (1 - X)/P(X = 0 | \mathbf{W})$, [Lemma 2](#) recovers the classic result that we can estimate the (unweighted) ATE by inverse probability weighting. If X is continuous with density $f_X(x | \mathbf{W})$ conditional on \mathbf{W} , letting $\alpha(x, \mathbf{W}) = -f'_X(x | \mathbf{W})/f_X(x | \mathbf{W})$ recovers the average derivative $E[g'(X, \mathbf{W})]$.

For both of these special cases, there is a wealth of papers studying how to best implement regression-based or weighting-based approaches to estimating $\theta(\omega)$, or combinations of both. Recent influential results in the cross-sectional literature (e.g., [Chernozhukov, Escanciano, Ichimura, Newey, and Robins, 2022](#); [Chernozhukov, Chetverikov, Demirer, Duflo, Hansen, Newey, and Robins, 2018](#)) highlight the advantages of combining both approaches using the Neyman orthogonal moment condition

$$\theta(\omega) = E[\mu(X, \mathbf{W}, g) + \alpha(X, \mathbf{W})(Y - g(X, \mathbf{W}))], \quad (30)$$

where $\mu(X, \mathbf{W}, g) = g(1, \mathbf{W}) - g(0, \mathbf{W})$ for the ATE and $\mu(X, \mathbf{W}, g) = g'(X, \mathbf{W})$ for the average derivative. This moment condition is orthogonal in the sense that it is insensitive to small perturbations in g , in contrast to the regression-based moment condition $\theta(\omega) = E[\mu(X, \mathbf{W}, g)]$. As a result, an orthogonal method-of-moments estimator based on (30) that plugs in first-stage estimates of g and α can be viewed as a debiased version of the plug-in estimator utilizing the regression-based moment condition. Actually, the moment condition (30) is not only orthogonal, but also *doubly robust*—insensitive to large perturbations in either g or α so that the orthogonal method-of-moments estimator remains consistent so long as any *one* of the first-stage estimators is consistent for α or g , even if the other estimator is inconsistent. For the binary treatment case, the orthogonal method-of-moments estimator corresponds to the classic augmented inverse probability weighted estimator of [Robins, Rotnitzky, and Zhao \(1994\)](#).

For i.i.d. data, it is popular to combine the orthogonal moment condition with cross-fitting ([Chernozhukov, Escanciano, Ichimura, Newey, and Robins, 2022](#); [Chernozhukov, Chetverikov, Demirer, Duflo, Hansen, Newey, and Robins, 2018](#)).²³ This allows for regularity conditions

²³Take a sample sum of the moment condition (30) over the first half of the sample, plugging in estimates $\hat{\alpha}_2$ and $\hat{\gamma}_2$ of α and γ constructed using the second half of the sample, and add to it a sample sum of the moment condition over the second half of the sample, but where we plug in estimates $\hat{\alpha}_1$ and $\hat{\gamma}_1$ from the *first* half of the sample.

that are weak enough to accommodate a variety of first-step estimators of α and γ , including kernels, series, as well as lasso, random forests, or other machine learning estimators, provided that these estimators converge sufficiently fast. Because of this flexibility, the approach is known as debiased machine learning. [Chernozhukov, Newey, and Singh \(2022\)](#) and [Hirshberg and Wager \(2021\)](#) develop alternatives to this approach that bypass the need to explicitly estimate the Riesz representer α .

These approaches all deliver estimators of $\theta(\omega)$ that converge, under appropriate regularity conditions, at the usual parametric rate (square root of sample size) even if the first-stage estimators are based on complicated nonparametric or machine learning algorithms. Recent work by [Gonçalves, Herrera, Kilian, and Pesavento \(2024a\)](#), [Ballinari and Wehrli \(2024\)](#), and [Ballarin \(2025\)](#) aims to adapt these semiparametric approaches to time series contexts. One potential worry is that even in the absence of covariates, given the small sample sizes typically available in macroeconomic applications, estimates of average marginal effects relying on machine-learning or nonparametric first-step estimates of the shock density and the structural function may yield estimates that are too noisy and sensitive to the choice of first-stage tuning parameters. When covariates are needed to argue that the observed variable X is exogenous, the data requirements become even more severe.

The practical challenges associated with fully nonparametric estimation motivates studying what the simple OLS local projection (7) estimates when the true regression function is nonlinear. Extending the analysis of [Section 3.3](#), we now allow the researcher to control for covariates more flexibly by considering the partially linear regression

$$Y = X\beta + \gamma(\mathbf{W}) + \text{residual}, \quad \text{where } \gamma \in \Gamma,$$

and Γ is a linear space of control functions that contains the constant function 1. This covers the case with a linear adjustment by letting $\Gamma = \{a + \mathbf{w}'\mathbf{b} : a \in \mathbb{R}, \mathbf{b} \in \mathbb{R}^{\dim(\mathbf{w})}\}$ be the class of linear functions of \mathbf{w} , as well as the semiparametric partially linear model that lets Γ be a large class of “nonparametric” functions. By the projection theorem, the estimand β in this regression is given by

$$\beta \equiv \frac{E[(X - \pi(\mathbf{W}))g(X, \mathbf{W})]}{\text{Var}(X - \pi(\mathbf{W}))}, \quad (31)$$

where $\pi(\mathbf{W}) \equiv \text{argmin}_{\gamma_0 \in \Gamma} E[(X - \gamma_0(\mathbf{W}))^2]$ denotes the projection of X onto Γ (if X and \mathbf{W} are independent, the estimand (31) reduces to that in (8); we assume the projection exists). We denote the true conditional expectation (the propensity score, if X is binary) by $\pi^*(\mathbf{W}) \equiv E[X \mid \mathbf{W}]$.

The next result uses [Lemma 2](#) to generalize [Proposition 1](#) to the case with covariates.

Proposition 7. *Let $\omega^*(x, \mathbf{W}) \equiv E[\mathbb{1}\{X \geq x\}(X - \pi^*(\mathbf{W})) \mid \mathbf{W}]$, and suppose X has finite second moments, and that $\text{Var}(X - \pi(\mathbf{W})) > 0$. Suppose that either (a) $\pi = \pi^*$; or else (b) for some $\gamma_0, \gamma_1 \in \Gamma$, and some weights $\lambda(x, \mathbf{w})$ such that $\int \lambda(x, \mathbf{w}) dx = \pi^*(\mathbf{w}) + \gamma_1(\mathbf{w})$, $E[g(X, \mathbf{W}) \mid \mathbf{W} = \mathbf{w}] = \gamma_0(\mathbf{w}) + \int \lambda(x, \mathbf{w}) g'(x, \mathbf{w}) dx$ for almost all \mathbf{w} .*

Furthermore, assume that conditional on \mathbf{W} , the following holds almost surely: (i) the support of X is contained in a (possibly unbounded) interval $I_{\mathbf{W}} \subseteq \mathbb{R}$; and (ii) $g(\cdot, \mathbf{W})$ is locally absolutely continuous on $I_{\mathbf{W}}$. Finally, assume that (iii) $E[\int |\omega^(x, \mathbf{W}) g'(x, \mathbf{W})| dx] < \infty$ and $E[|g(X, \mathbf{W})(X - \pi^*(\mathbf{W}))|] < \infty$. Then the estimand [\(31\)](#) satisfies*

$$\beta = \theta(\omega), \quad \text{where} \quad \omega(x, \mathbf{W}) \equiv \frac{\omega^*(x, \mathbf{W}) + (\pi^*(\mathbf{W}) - \pi(\mathbf{W}))\lambda(x, \mathbf{W})}{\text{Var}(X - \pi(\mathbf{W}))},$$

where, if condition (a) holds, we let $\lambda(x, \mathbf{w}) = 0$.

The weights integrate to one: $E[\int \omega(x, \mathbf{W}) dx] = 1$. A sufficient condition for the weights to be non-negative is that condition (a) holds, in which case $\omega^(x, \mathbf{w})$ is hump-shaped as a function of x for almost all \mathbf{w} : monotonically increasing from 0 to its maximum for $x = \pi^*(\mathbf{w})$, and then monotonically decreasing back to 0.*

To interpret this result, it is useful to first consider the case where the class Γ of control functions is rich enough so that condition (a) holds. This is trivially the case, for instance, if \mathbf{W} just consists of a single set of fixed effects. In this case, we obtain an analogue of [Proposition 1](#): (partially) linear regression identifies a weighted average of marginal effects, with hump-shaped weights—this holds regardless of the distribution of X , be it discrete, continuous or mixed. If the conditional distribution of X given \mathbf{W} is continuous, then the weighting $\omega(x, \mathbf{W}) = \omega^*(x, \mathbf{W}) / \text{Var}(X - \pi^*(\mathbf{W}))$ varies smoothly with x ; if there are mass points, as in the case of a discrete or mixed distribution, then the weight function jumps discontinuously at the mass points. If the treatment X is discrete with support $0, 1, \dots$, we recover the result in [Angrist and Krueger \(1999\)](#) that regression estimates a weighted average of the causal effects of increasing X by one unit,

$$\beta = \frac{E[\sum_{s=1}^{\infty} E[\mathbb{1}\{X \geq s\}(X - \pi^*(\mathbf{W})) \mid \mathbf{W}](g(s, \mathbf{W}) - g(s-1, \mathbf{W}))]}{E[\sum_{s=1}^{\infty} E[\mathbb{1}\{X \geq s\}(X - \pi^*(\mathbf{W})) \mid \mathbf{W}]]}.$$

Now suppose that condition (a) is violated, because the specification for Γ is not sufficiently flexible to model the true propensity score $\pi^*(\mathbf{W})$. In general, this may lead to

omitted variable bias, as the partially linear model may not be sufficiently flexible to account for all confounding due to \mathbf{W} . Condition (b) prevents this scenario, ensuring that any bias due to confounding is accounted for. As a simple example of when the condition holds, consider the case where the true conditional mean function has a multiplicative form: $g(X, \mathbf{W}) = Xg'(\mathbf{W}) + \gamma_0(\mathbf{W})$, with $\gamma_0 \in \Gamma$. The partially linear model is misspecified, because the marginal effect is not constant, but varies with \mathbf{W} . But condition (b) holds with $\gamma_1(\mathbf{w}) = 0$ and $\lambda(x, \mathbf{w}) = \pi^*(\mathbf{w})\varphi(x)$, where $\varphi(x)$ is an arbitrary density function. Because the marginal effect varies only with \mathbf{W} but not with X , [Proposition 7](#) simplifies to

$$\beta = \frac{E[(E[X^2 | \mathbf{W}] - \pi(\mathbf{W})\pi^*(\mathbf{W}))g'(\mathbf{W})]}{\text{Var}(X - \pi^*(\mathbf{W}))}.$$

If $\pi = \pi^*$, we obtain a generalization of the [Angrist \(1998\)](#) result for binary treatments: the weights are proportional to the conditional variance of X , $\text{Var}(X | \mathbf{W})$. But if $\pi \neq \pi^*$, the weight function may be negative for some values of \mathbf{W} . Consider, for instance, a panel data scenario where Γ is linear, and \mathbf{W} consists of unit and time fixed effects. Then the assumption $\gamma_0 \in \Gamma$ amounts to a parallel trends assumption: in the absence of the treatment, the average differences in outcomes for different units are constant and do not depend on the time period. If the treatment effects are heterogeneous, so that $g'(\mathbf{W})$ depends on the unit and time period, then this two-way fixed effects regression still estimates a weighted average of marginal effects, but with weights that are negative if $E[X^2 | \mathbf{W}] < \pi(\mathbf{W})\pi^*(\mathbf{W})$. In the context of a binary treatment and two-way fixed effects regressions, this result has been noted in [de Chaisemartin and D'Haultfœuille \(2020\)](#) and [Goodman-Bacon \(2021\)](#). [Proposition 7](#) generalizes this to an arbitrary treatment distribution and a general regression specification.

Another example where condition (b) of [Proposition 7](#) holds is when X is bounded below by some baseline value (say, 0), $X \geq 0$, and the baseline outcome model is correctly specified, $g(0, \mathbf{W}) \in \Gamma$. Then condition (b) holds with $\gamma_0(\mathbf{W}) = g(0, \mathbf{W})$, $\gamma_1 = 0$, and $\lambda(x, \mathbf{W}) = P(X \geq x | \mathbf{W})$. In this case, the weights simplify to $\omega(x, \mathbf{W}) = E[\mathbb{1}\{X \geq x\}(X - \pi(\mathbf{W})) | \mathbf{W}] / \text{Var}(X - \pi(\mathbf{W}))$.

Thus, if the marginal effects are constant, the partially linear model is doubly robust: the regression estimand is consistent for this constant treatment effect so long as either $\pi^* \in \Gamma$ or $g(0, \mathbf{W}) \in \Gamma$, as noted, for instance, in [Robins, Mark, and Newey \(1992\)](#). But this double robustness doesn't fully extend to the case with heterogeneous marginal effects. If the researcher gets the model for X right, in the sense that $\pi = \pi^*$, then the partially linear regression estimates an average marginal effect. However, if the researcher gets it wrong,

and only gets the outcome model under no treatment right, so that only condition (b) of [Proposition 7](#) holds, then weights on some of the true marginal effects $g'(X, \mathbf{W})$ may be negative, risking a sign-reversal. This asymmetry has been noted by [Goldsmith-Pinkham, Hull, and Kolesár \(2024\)](#) for the case with a binary treatment X .²⁴ [Proposition 7](#) shows that the result is general. The upshot of this asymmetry is that in cases where the treatment variable X is only conditionally exogenous—whether the data is cross-sectional, panel, or time series—it pays off to conduct a sensitivity analysis with respect to the functional form of the control specification, in addition to the standard sensitivity analysis with respect to the set of controls.

7 Conclusion

We have shown that conventional linear methods for identifying causal effects in applied time series analysis based on observed shocks or proxies are robust to misspecification: they estimate a positively weighted average of the true nonlinear causal effects, irrespective of the extent of nonlinearities in the underlying DGP. By contrast, identification approaches that exploit heteroskedasticity or non-Gaussianity of latent shocks are highly sensitive to violations of the assumed linear functional form of the structural model. Moreover, while linear identification via heteroskedasticity provides some testable restrictions, identification via non-Gaussianity is generally unfalsifiable despite the potential for severe biases.

Our results suggest that it is worthwhile for applied researchers who are interested in average marginal effects to expend the effort involved in constructing direct measures of shocks, or at least proxies that are credibly (approximately) monotonically related to the latent shock of interest. Shock measurement via narrative approaches or detailed institutional knowledge is admittedly highly work-intensive and in many applications may be practically impossible. Nevertheless, when the strategy is feasible, it affords an insurance against functional form misspecification which is not matched by the other identification approaches that we analyze. This is similar to the robustness of linear regression for estimating average treatment effects in randomized control trials relative to estimators that rely on parametric adjustment for nonlinearities or selection in observational data. Our results do not directly speak to other identification approaches like identification via long-run or sign restrictions, and we leave

²⁴[Goldsmith-Pinkham, Hull, and Kolesár \(2024\)](#) also show that in regressions on multiple mutually exclusive treatment indicators, the regression estimand on a given treatment contains an additional contamination bias term corresponding to non-convex average effects of the other treatments.

these for future work.²⁵

When reporting impulse responses from linear specifications with observed shocks, we recommend that researchers routinely report the implicit weight function, which is easy to compute via standard regression software. If the weight function associated with conventional local projections or VARs is deemed to be unattractive, one can target other causal summaries as discussed in [Section 6](#), though further analysis is required on the best practices for doing so in macroeconomic applications.

More broadly, we hope that our paper will boost the agenda spearheaded by [White \(2006\)](#), [White and Kennedy \(2009\)](#), [Angrist and Kuersteiner \(2011\)](#), [Angrist, Jordà, and Kuersteiner \(2018\)](#), [Rambachan and Shephard \(2021\)](#), and [Gonçalves, Herrera, Kilian, and Pesavento \(2024a,b\)](#) that seeks to draw lessons for macroeconometrics from the microeconomic treatment effect literature. The nonparametric framework used in the treatment effect literature contains useful lessons for empirical work in macroeconomics, even though conventional approaches to nonparametric estimation or debiased machine learning methods are impractical due to the much smaller data sets typical in macroeconomics.

²⁵As a referee pointed out to us, it may be possible to analyze long-run restrictions by extending [Proposition 3](#) to apply to the instrumental variable specification in [Shapiro and Watson \(1988\)](#). However, since the derivations in the latter paper rely on linearity of the structural function, it is not immediately clear what assumptions would guarantee convex weights.

References

- ANGRIST, J. D. (1998): “Estimating the Labor Market Impact of Voluntary Military Service Using Social Security Data on Military Applicants,” *Econometrica*, 66(2), 249–288.
- (2001): “Dummy Endogenous Regressors: Simple Strategies for Empirical Practice,” *Journal of Business & Economic Statistics*, 19(1), 2–16.
- ANGRIST, J. D., K. GRADDY, AND G. W. IMBENS (2000): “The Interpretation of Instrumental Variables Estimators in Simultaneous Equations Models with an Application to the Demand for Fish,” *Review of Economic Studies*, 67(3), 499–527.
- ANGRIST, J. D., Ò. JORDÀ, AND G. M. KUERSTEINER (2018): “Semiparametric Estimates of Monetary Policy Effects: String Theory Revisited,” *Journal of Business & Economic Statistics*, 36(3), 371–387.
- ANGRIST, J. D., AND A. B. KRUEGER (1999): “Empirical Strategies in Labor Economics,” in *Handbook of Labor Economics*, ed. by O. C. Ashenfelter, and D. Card, vol. 3, chap. 3, pp. 1277–1366. Elsevier.
- ANGRIST, J. D., AND G. M. KUERSTEINER (2011): “Causal Effects of Monetary Shocks: Semiparametric Conditional Independence Tests with a Multinomial Propensity Score,” *Review of Economics and Statistics*, 93(3), 725–747.
- ANGRIST, J. D., AND J.-S. PISCHKE (2009): *Mostly Harmless Econometrics*. Princeton University Press.
- ANTOLÍN-DÍAZ, J., AND J. F. RUBIO-RAMÍREZ (2018): “Narrative Sign Restrictions for SVARs,” *American Economic Review*, 108(10), 2802–29.
- BACCHIOCCHI, E., A. BASTIANIN, T. KITAGAWA, AND E. MIRTO (2024): “Partially identified heteroskedastic SVARs,” arXiv:2403.06879.
- BALLARIN, G. (2025): “Impulse Response Analysis of Structural Nonlinear Time Series Models,” arXiv:2305.19089.
- BALLINARI, D., AND A. WEHRLI (2024): “Semiparametric Inference for Impulse Response Functions Using Double/Debiased Machine Learning,” arXiv:2411.10009.

- BEN ZEEV, N., AND E. PAPP (2017): “Chronicle of A War Foretold: The Macroeconomic Effects of Anticipated Defence Spending Shocks,” *The Economic Journal*, 127(603), 1568–1597.
- BLANCHARD, O., AND R. PEROTTI (2002): “An Empirical Characterization of the Dynamic Effects of Changes in Government Spending and Taxes on Output,” *The Quarterly Journal of Economics*, 117(4), 1329–1368.
- BLANDHOL, C., J. BONNEY, M. MOGSTAD, AND A. TORGOVITSKY (2022): “When is TSLS Actually LATE?,” Working Paper 2022-16, Becker Friedman Institute for Economics, University of Chicago.
- CARAVELLO, T. E., AND P. MARTÍNEZ BRUERA (2024): “Disentangling sign and size non-linearities,” Manuscript, MIT.
- CARROLL, R. J., D. RUPPERT, L. A. STEFANSKI, AND C. M. CRAINICEANU (2006): *Measurement Error in Nonlinear Models*. Chapman and Hall/CRC, New York, NY, 2nd edn.
- CASINI, A., AND A. MCCLOSKEY (2024): “Identification and Estimation of Causal Effects in High-Frequency Event Studies,” Manuscript, University of Rome Tor Vergata.
- CHERNOZHUKOV, V., D. CHETVERIKOV, M. DEMIRER, E. DUFLO, C. HANSEN, W. NEWEY, AND J. M. ROBINS (2018): “Double/Debiased Machine Learning for Treatment and Structural Parameters,” *The Econometrics Journal*, 21(1), C1–C68.
- CHERNOZHUKOV, V., J. C. ESCANCIANO, H. ICHIMURA, W. K. NEWEY, AND J. M. ROBINS (2022): “Locally Robust Semiparametric Estimation,” *Econometrica*, 90(4), 1501–1535.
- CHERNOZHUKOV, V., W. K. NEWEY, AND R. SINGH (2022): “Automatic Debiased Machine Learning of Causal and Structural Effects,” *Econometrica*, 90(3), 967–1027.
- CHRISTIANO, L., M. EICHENBAUM, AND C. EVANS (1999): “Monetary Policy Shocks: What Have We Learned and to What End?,” in *Handbook of Macroeconomics*, ed. by J. B. Taylor, and M. Woodford, vol. 1A, pp. 65–148. Elsevier.
- CLOYNE, J., O. JORDÀ, AND A. M. TAYLOR (2023): “State-Dependent Local Projections: Understanding Impulse Response Heterogeneity,” Working Paper 30971, National Bureau of Economic Research.

- COMON, P. (1994): “Independent component analysis, a new concept?,” *Signal Processing*, 36(3), 287–314.
- DE CHAISEMARTIN, C., AND X. D’HAULTFŒUILLE (2020): “Two-Way Fixed Effects Estimators with Heterogeneous Treatment Effects,” *American Economic Review*, 110(9), 2964–2996.
- DUDLEY, R. M. (2002): *Real Analysis and Probability*. Cambridge University Press, Cambridge.
- FAN, J., AND Q. YAO (2003): *Nonlinear Time Series: Nonparametric and Parametric Methods*, Springer Series in Statistics. Springer.
- FISHER, J. D., AND R. PETERS (2010): “Using Stock Returns to Identify Government Spending Shocks,” *The Economic Journal*, 120(544), 414–436.
- GALLANT, A. R., P. E. ROSSI, AND G. TAUCHEN (1993): “Nonlinear Dynamic Structures,” *Econometrica*, 61(4), 871–871.
- GIACOMINI, R., T. KITAGAWA, AND M. READ (2023): “Identification and Inference under Narrative Restrictions,” Research Discussion Paper 2023-07, Reserve Bank of Australia.
- GOLDSMITH-PINKHAM, P., P. HULL, AND M. KOLESÁR (2024): “Contamination Bias in Linear Regressions,” *American Economic Review*, 114(12), 4015–51.
- GONÇALVES, S., A. M. HERRERA, L. KILIAN, AND E. PESAVENTO (2021): “Impulse response analysis for structural dynamic models with nonlinear regressors,” *Journal of Econometrics*, 225(1), 107–130, Themed Issue: “Vector Autoregressions”.
- (2024a): “Nonparametric Local Projections,” Working Paper 2414, Federal Reserve Bank of Dallas.
- (2024b): “State-dependent local projections,” *Journal of Econometrics*, 244(2), 105702.
- GOODMAN-BACON, A. (2021): “Difference-in-Differences with Variation in Treatment Timing,” *Journal of Econometrics*, 225(2), 254–277.

- GOURIÉROUX, C., AND Q. LEE (2023): “Nonlinear Impulse Response Functions and Local Projections,” Manuscript, University of Toronto.
- GOURIÉROUX, C., A. MONFORT, AND J.-P. RENNE (2017): “Statistical inference for independent component analysis: Application to structural VAR models,” *Journal of Econometrics*, 196(1), 111–126.
- GRAHAM, B. S., AND C. C. DE XAVIER PINTO (2022): “Semiparametrically efficient estimation of the average linear regression function,” *Journal of Econometrics*, 226(1), 115–138, Annals Issue in Honor of Gary Chamberlain.
- GUNSILIUS, F., AND S. SCHENNACH (2023): “Independent Nonlinear Component Analysis,” *Journal of the American Statistical Association*, 118(542), 1305–1318.
- HÄRDLE, W., AND T. M. STOKER (1989): “Investigating Smooth Multiple Regression by the Method of Average Derivatives,” *Journal of the American Statistical Association*, 84(408), 986–995.
- HERBST, E. P., AND B. K. JOHANSEN (2024): “Bias in local projections,” *Journal of Econometrics*, 240(1), 105655.
- HERBST, E. P., AND F. SCHORFHEIDE (2016): *Bayesian Estimation of DSGE Models*. Princeton University Press, Princeton.
- HIRSHBERG, D. A., AND S. WAGER (2021): “Augmented Minimax Linear Estimation,” *The Annals of Statistics*, 49(6), 3206–3227.
- HYVÄRINEN, A., J. KARHUNEN, AND E. OJA (2001): *Independent Component Analysis*. John Wiley & Sons, Inc.
- IMBENS, G. W., AND J. D. ANGRIST (1994): “Identification and Estimation of Local Average Treatment Effects,” *Econometrica*, 62(2), 467–475.
- JORDÀ, Ò. (2005): “Estimation and Inference of Impulse Responses by Local Projections,” *American Economic Review*, 95(1), 161–182.
- (2023): “Local Projections for Applied Economics,” *Annual Review of Economics*, 15, 607–631.

- JORDÀ, O., AND A. M. TAYLOR (2025): “Local Projections,” *Journal of Economic Literature*, 63(1), 59–110.
- JUTTEN, C., AND J. KARHUNEN (2003): “Advances in Nonlinear Blind Source Separation,” in *4th International Symposium on Independent Component Analysis and Blind Signal Separation*, Nara, Japan.
- KAGAN, A., J. LINNIK, AND C. RAO (1973): *Characterization Problems in Mathematical Statistics*. Wiley.
- KILIAN, L., AND H. LÜTKEPOHL (2017): *Structural Vector Autoregressive Analysis*. Cambridge University Press.
- KITAGAWA, T., W. WANG, AND M. XU (2023): “Policy Choice in Time Series by Empirical Welfare Maximization,” arXiv:2205.03970.
- KLEIN, R., AND F. VELLA (2010): “Estimating a class of triangular simultaneous equations models without exclusion restrictions,” *Journal of Econometrics*, 154(2), 154–164.
- LANNE, M., M. MEITZ, AND P. SAIKKONEN (2017): “Identification and estimation of non-Gaussian structural vector autoregressions,” *Journal of Econometrics*, 196(2), 288–304.
- LEWBEL, A. (2012): “Using Heteroscedasticity to Identify and Estimate Mismeasured and Endogenous Regressor Models,” *Journal of Business & Economic Statistics*, 30(1), 67–80.
- (2018): “Identification and Estimation Using Heteroscedasticity without Instruments: The Binary Endogenous Regressor Case,” *Economics Letters*, 165, 10–12.
- LEWBEL, A., S. M. SCHENNACH, AND L. ZHANG (2024): “Identification of a Triangular Two Equation System Without Instruments,” *Journal of Business & Economic Statistics*, 42(1), 14–25.
- LEWIS, D. J. (2024): “Identification Based on Higher Moments,” Manuscript, UCL.
- McKAY, A., AND C. K. WOLF (2023): “What Can Time-Series Regressions Tell Us About Policy Counterfactuals?,” *Econometrica*, 91(5), 1695–1725.
- MONTIEL OLEA, J. L., M. PLAGBORG-MØLLER, AND E. QIAN (2022): “SVAR Identification from Higher Moments: Has the Simultaneous Causality Problem Been Solved?,” *AEA Papers and Proceedings*, 112, 481–485.

- NEWHEY, W. K., AND T. M. STOKER (1993): “Efficiency of Weighted Average Derivative Estimators and Index Models,” *Econometrica*, 61(5), 1199–1223.
- PLAGBORG-MØLLER, M., AND C. K. WOLF (2021): “Local Projections and VARs Estimate the Same Impulse Responses,” *Econometrica*, 89(2), 955–980.
- PLAGBORG-MØLLER, M. (2022): “Discussion of “Narrative Restrictions and Proxies” by Raffaella Giacomini, Toru Kitagawa, and Matthew Read,” *Journal of Business & Economic Statistics*, 40(4), 1434–1437.
- POTTER, S. M. (2000): “Nonlinear impulse response functions,” *Journal of Economic Dynamics and Control*, 24(10), 1425–1446.
- POWELL, J. L., J. H. STOCK, AND T. M. STOKER (1989): “Semiparametric Estimation of Index Coefficients,” *Econometrica*, 57(6), 1403–1430.
- RAMBACHAN, A., AND N. SHEPHARD (2021): “When do common time series estimands have nonparametric causal meaning?,” Manuscript, Harvard University.
- RAMEY, V. A. (2011): “Identifying Government Spending Shocks: It’s all in the Timing,” *The Quarterly Journal of Economics*, 126(1), 1–50.
- (2016): “Macroeconomic Shocks and Their Propagation,” in *Handbook of Macroeconomics*, ed. by J. B. Taylor, and H. Uhlig, vol. 2, chap. 2, pp. 71–162. Elsevier.
- RIGOBON, R. (2003): “Identification Through Heteroskedasticity,” *Review of Economics and Statistics*, 85(4), 777–792.
- RIGOBON, R., AND B. SACK (2004): “The impact of monetary policy on asset prices,” *Journal of Monetary Economics*, 51(8), 1553–1575.
- ROBINS, J., M. SUED, Q. LEI-GOMEZ, AND A. ROTNITZKY (2007): “Comment: Performance of Double-Robust Estimators When “Inverse Probability” Weights Are Highly Variable,” *Statistical Science*, 22(4), 544–559.
- ROBINS, J. M., S. D. MARK, AND W. K. NEWHEY (1992): “Estimating Exposure Effects by Modelling the Expectation of Exposure Conditional on Confounders,” *Biometrics*, 48(2), 479–495.

- ROBINS, J. M., A. ROTNITZKY, AND L. P. ZHAO (1994): “Estimation of Regression Coefficients When Some Regressors Are Not Always Observed,” *Journal of the American Statistical Association*, 89(427), 846–866.
- SENTANA, E., AND G. FIORENTINI (2001): “Identification, estimation and testing of conditionally heteroskedastic factor models,” *Journal of Econometrics*, 102(2), 143–164.
- SHAPIRO, M. D., AND M. W. WATSON (1988): “Sources of Business Cycle Fluctuations,” *NBER Macroeconomics Annual*, 3, 111–148.
- SMALL, D. S., Z. TAN, R. R. RAMSAHAI, S. A. LORCH, AND M. A. BROOKHART (2017): “Instrumental Variable Estimation with a Stochastic Monotonicity Assumption,” *Statistical Science*, 32(4), 561–579.
- STEIN, C. M. (1981): “Estimation of the Mean of a Multivariate Normal Distribution,” *Annals of Statistics*, 9(6), 1135–1151.
- STOCK, J. H., AND M. W. WATSON (2018): “Identification and Estimation of Dynamic Causal Effects in Macroeconomics Using External Instruments,” *The Economic Journal*, 128(610), 917–948.
- WHITE, H. (1980): “Using Least Squares to Approximate Unknown Regression Functions,” *International Economic Review*, 21(1), 149–170.
- (2006): “Time-series estimation of the effects of natural experiments,” *Journal of Econometrics*, 135(1-2), 527–566.
- WHITE, H., AND P. KENNEDY (2009): “Retrospective Estimation of Causal Effects Through Time,” in *The Methodology and Practice of Econometrics: A Festschrift in Honour of David F. Hendry*, ed. by J. Castle, and N. Shephard, chap. 3, pp. 59–87. Oxford University Press.
- WRIGHT, J. H. (2012): “What does Monetary Policy do to Long-term Interest Rates at the Zero Lower Bound?,” *The Economic Journal*, 122(564), F447–F466.
- YITZHAKI, S. (1996): “On Using Linear Regressions in Welfare Economics,” *Journal of Business & Economic Statistics*, 14(4), 10.